



sfl[phone]

The free software
enterprise-class
softphone for GNU/Linux



debian



kubuntu



openSUSE



Não profissional





OPOVO



Novell.
PartnerNet®
DATA CENTER
SPECIALIST



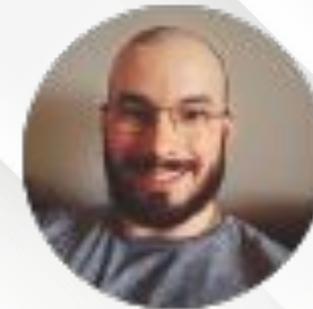
flipp



**Certified
LINUX
ADMINISTRATOR**
Novell.



Marcelo Cavalcante Rocha



 marcelo@marcelocavalcante.net

www.marcelocavalcante.net

 www.linkedin.com/in/marcelocrocha

 marcelokalib

 marcelocavalcantecrocha

 facebook.com/marcelocavalcante

kalib @ freenode.org





Kubernetes

Conceitos - Core

Arquitetura do Cluster

Um pouco sobre a arquitetura de um cluster Kubernetes e seus componentes básicos e fundamentais

Serviços e outras primitivas de Redes

Introdução a como todos esses componentes se comunicam bem como sobre a estrutura de redes do Kubernetes

APIs Primitivas

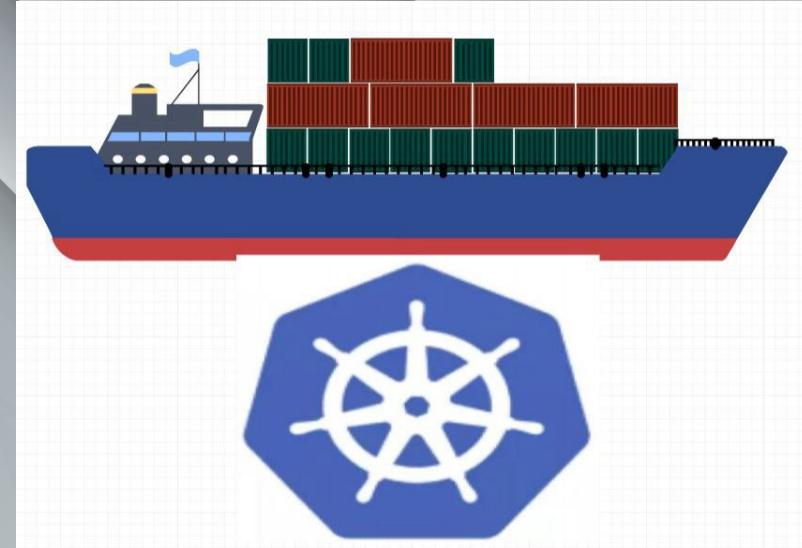
Introdução as APIs fundamentais do Kubernetes como pods, replicaset, deployments, services...

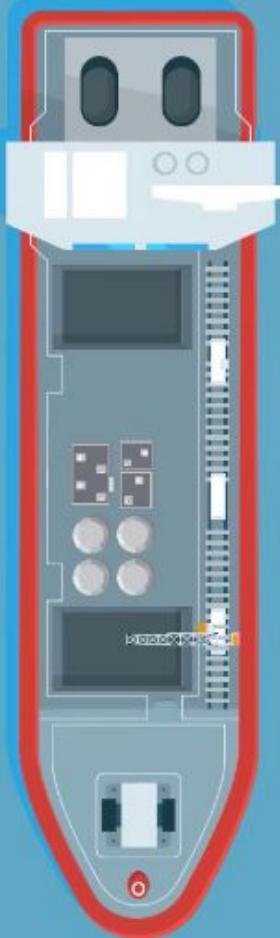


Arquitetura do Cluster

- △ Arquitetura
- △ ETCD
- △ Kube-API Server
- △ Controller Managers
- △ Kube Scheduler
- △ Kubelet
- △ Kube Proxy







MASTER NODE

Gerencia, Planeja, Agenda e Monitora os nodes

WORKER NODES

Hospeda aplicações de containers

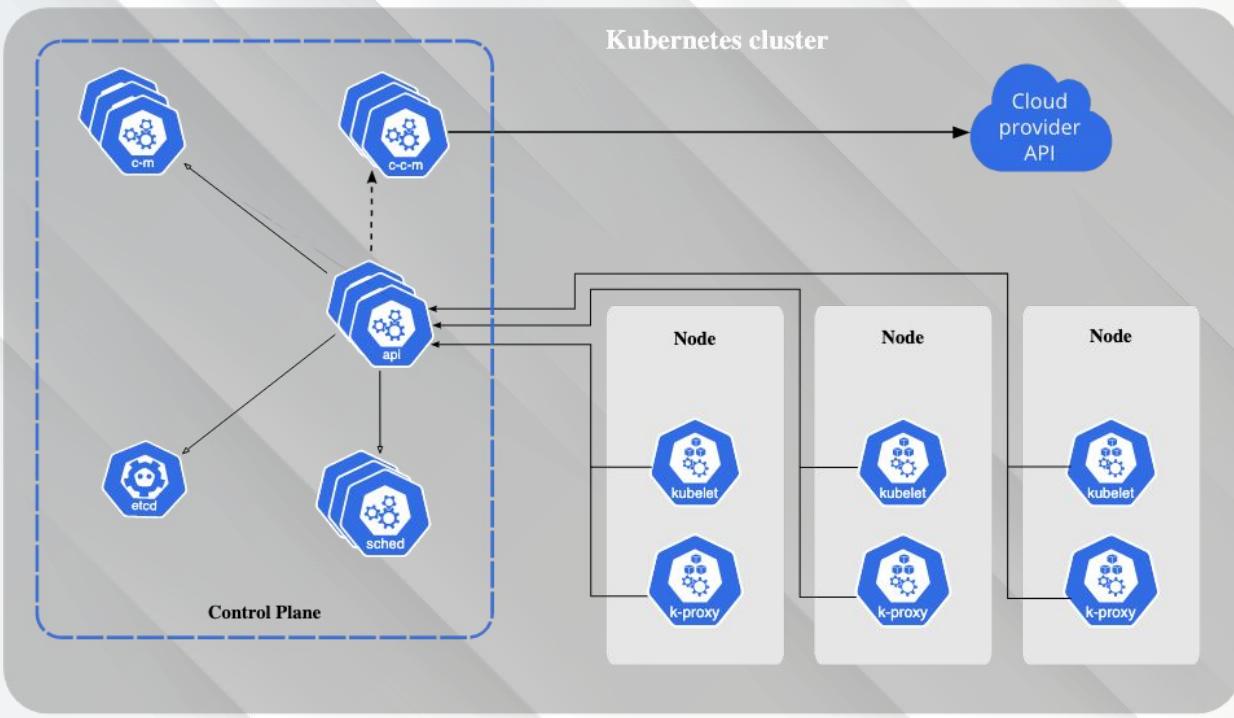
Um **cluster Kubernetes** é composto por **NODES** (ou nós), responsáveis por hospedar aplicações em formato de containers.

São como **navios de carga**, que apenas recebem os containers e os transportam. Na arquitetura de um cluster Kubernetes, chamamos estes nodes de **WORKER NODES**.

Os **navios de controle** são chamados de **MASTER NODE**. Eles são os responsáveis por monitorar e gerenciar o trabalho dos worker nodes: alocar containers, registrar informações, escolher em qual worker node vai cada container, etc.

<https://kubernetes.io/docs/concepts/overview/components/>

Control Plane Components





Master
Gerencia, Planeja, Agenda e
Monitora os nodes

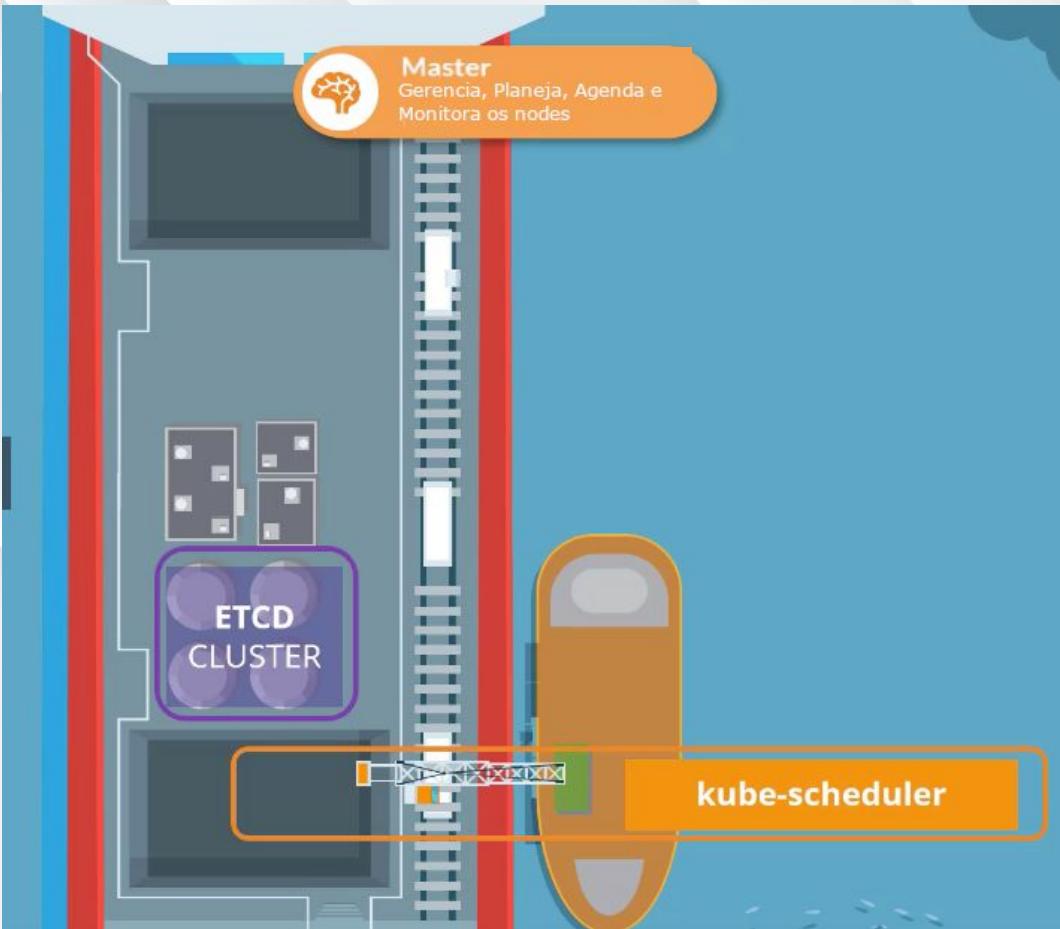


ETCD CLUSTER



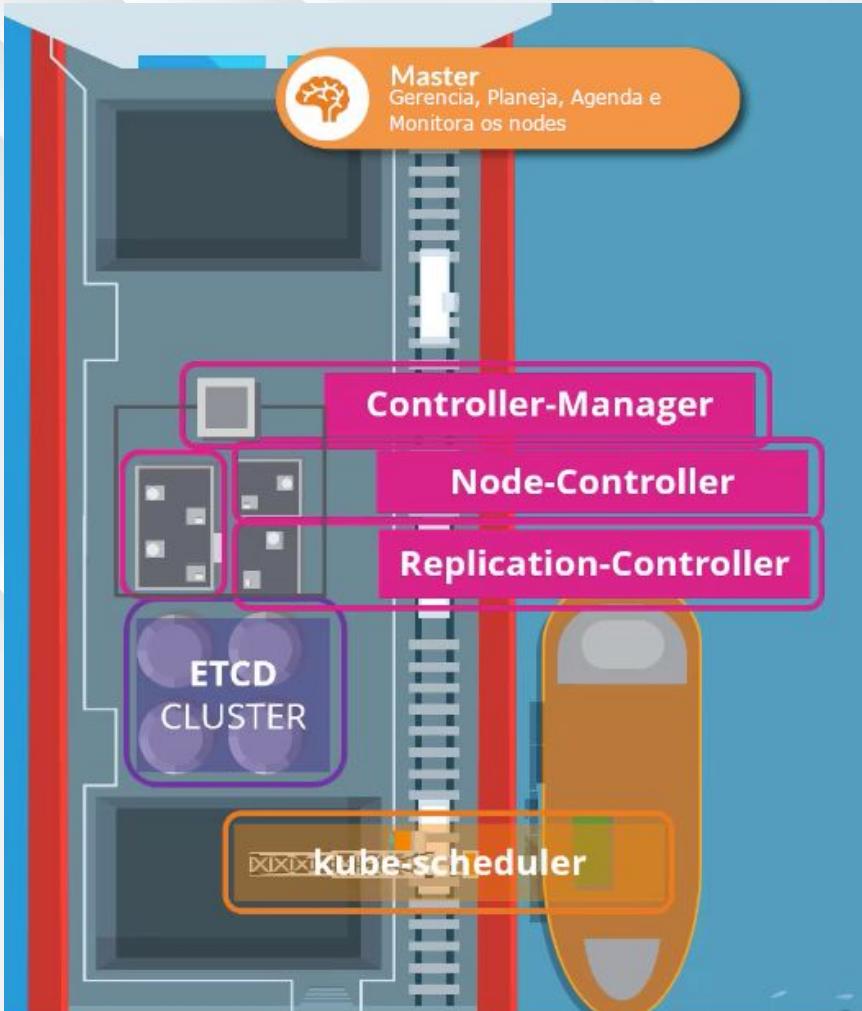
O Kubernetes utiliza-se do **ETCD** para armazenar todas as informações sobre os containers e nodes existentes no cluster.

É uma popular ferramenta de bancos de dados que armazena informações no formato **chave:valor**.



O componente **Scheduler** ou **Kube-Scheduler** é responsável por alocar os containers nos worker nodes apropriados de acordo com uma série de fatores como capacidade do worker node, tamanho do container, destino, tipo de aplicação, exigências e restrições do container ou worker node, etc.

Ele é basicamente o guindaste ou escritório de guindaste, que coloca os containers em cada navio de carga.



Controller-Managers são responsáveis por controles específicos no cluster.

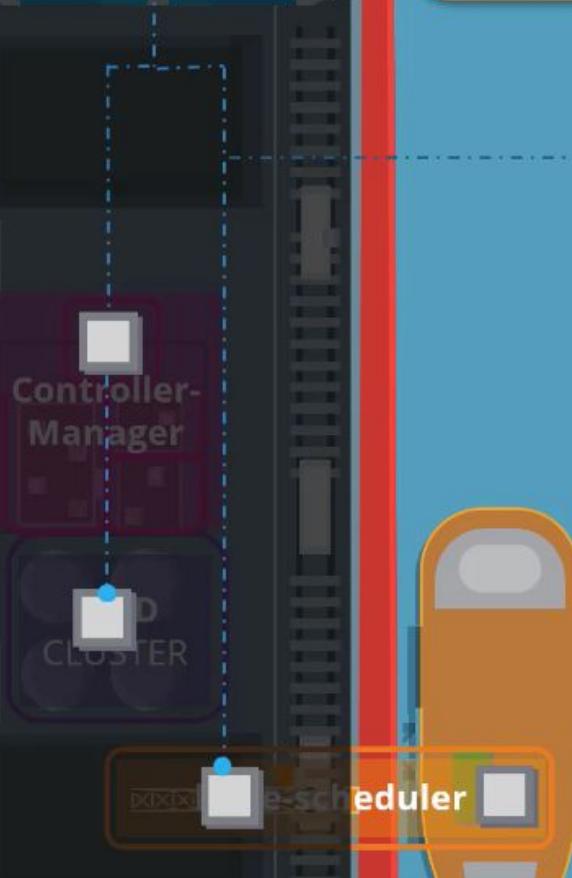
Exemplos de controllers são:

- **Node-Controller**: responsável pelos **nodes** em geral
- **Replication-Controller**: responsável pelas **réplicas**, ou seja o número de containers rodando em um determinado grupo de réplicas

kube-apiserver



Master
Gerencia, Planeja, Agenda e
Monitora os nodes



KUBE-APISERVER

O gerente entre os gerentes.

Responsável por **orquestrar** TODAS as operações no cluster.

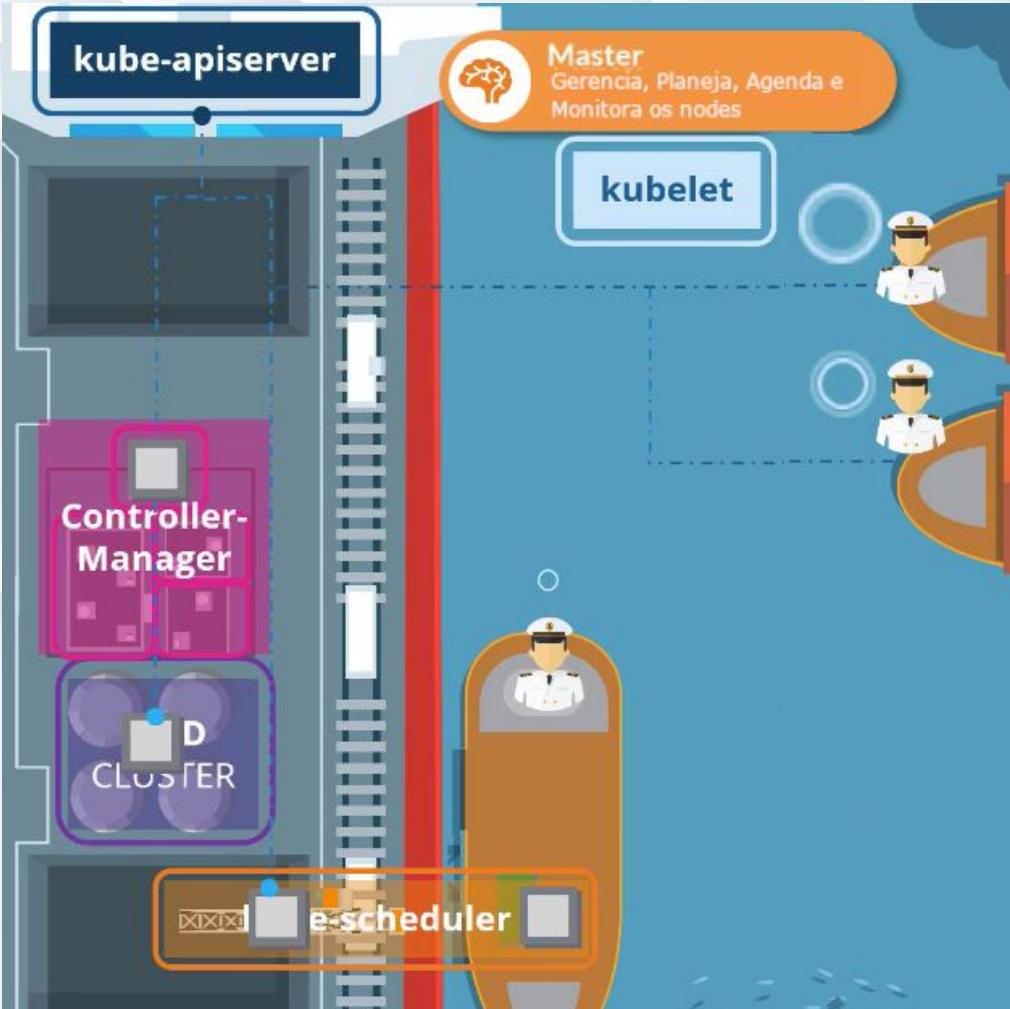
Expõe as **APIs do kubernetes** para que sejam utilizadas por usuários externos durante suas requisições e operações no cluster.

Assim como os demais componentes do Kubernetes, o **kube-apiserver** nada mais é do que um **container** no cluster.

CONTAINERS, CONTAINERS

EVERYWHERE

memegenerator.net

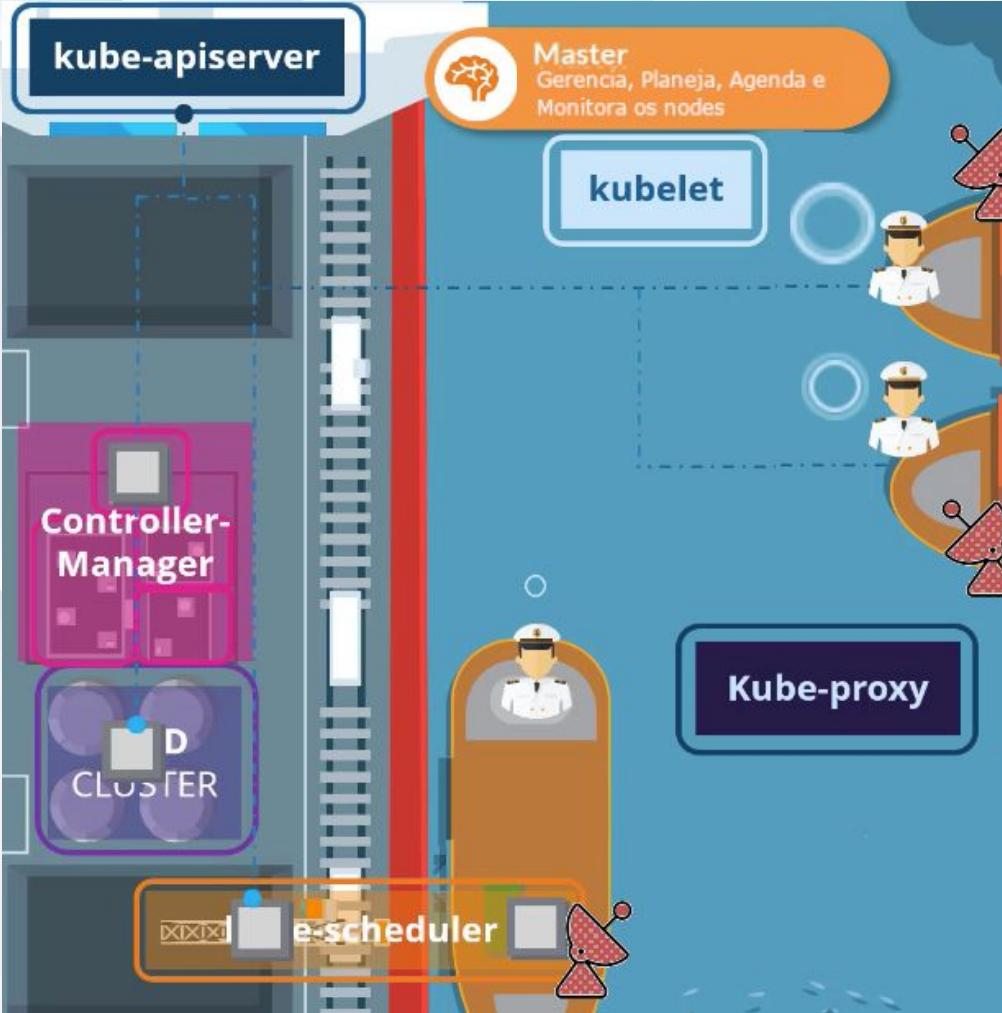


Kubelet

Todo navio precisa de um capitão e este capitão é o **Kubelet**.

Responsável pela **comunicação** entre o **worker node** e o **master**. Se **reporta** ao kube-apiserver constantemente passando informações sobre o **status** do seu node bem como dos containers que ali estão.

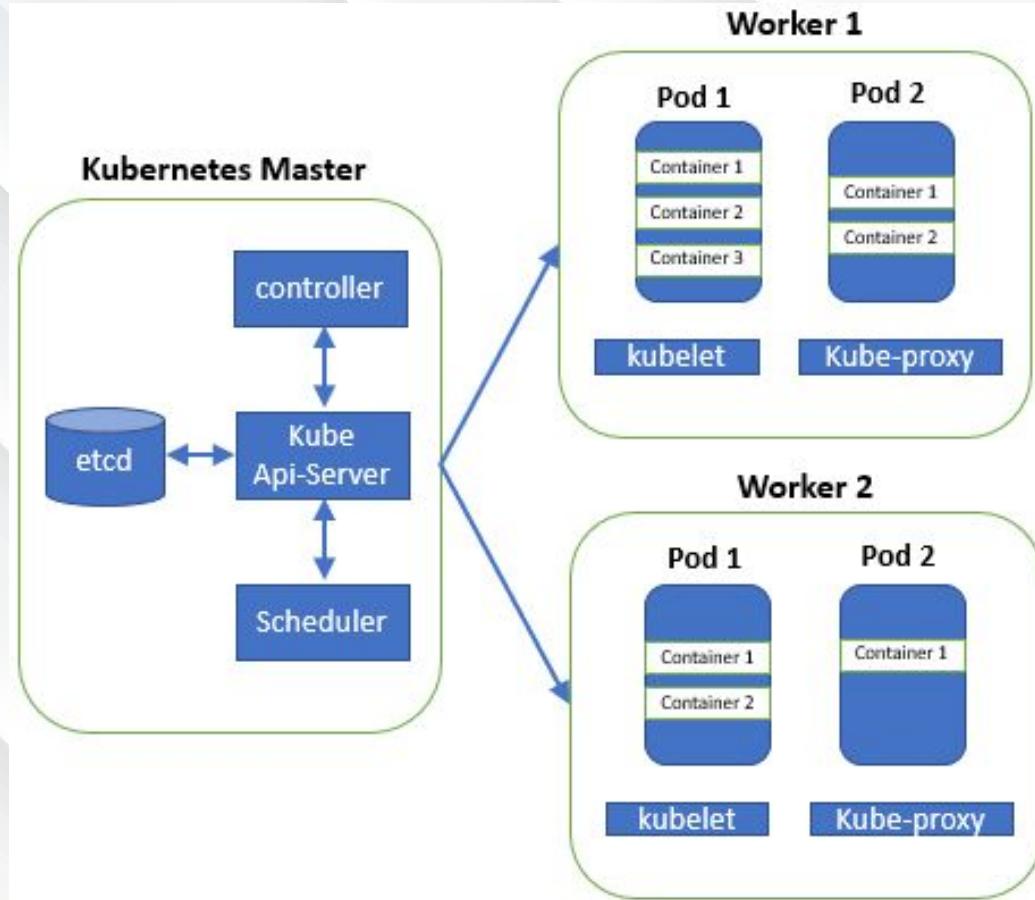
Ele é quem recebe as **instruções** do kube-apiserver sobre o que deve ser feito no seu node, containers que devem ser carregados lá, etc.



Kube-proxy

O kube-proxy é o serviço responsável por garantir a comunicação entre os containers/aplicações existentes no cluster de nodes seguindo uma série de regras para isso.

Resumão da arquitetura básica do Kubernetes





- O que é o ETCD?
- O que é um dado de formato Chave-Valor?
- Como começar?
- Como operar o ETCD?
- O que é um sistema distribuído?
- Como o ETCD opera?
- Protocolo RAFT
- Melhores práticas em números de nodes

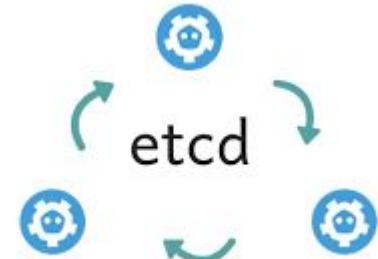
- Armazenamento de dados no formato chave-valor confiável, distribuído, e que é simples, seguro e rápido.

Key-Value (Chave-Valor)

Chave	Valor
Nome	Clevin
Idade	78
Localização	Paris

Tabular - Bancos de Dados Relacionais

Nome	Idade	Localização
Berg	18	Sobral
Bruno	45	Cabrobró
Samuel	32	Salvador



- Armazenamento de dados no formato chave-valor confiável, distribuído, e que é simples, seguro e rápido.

Key-Value (Chave-Valor)

Chave	Valor
Nome	Clevin
Idade	78
Localização	Paris

Put Nome "Clevin"

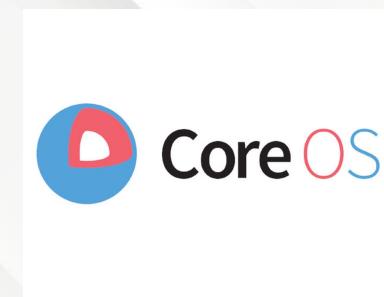
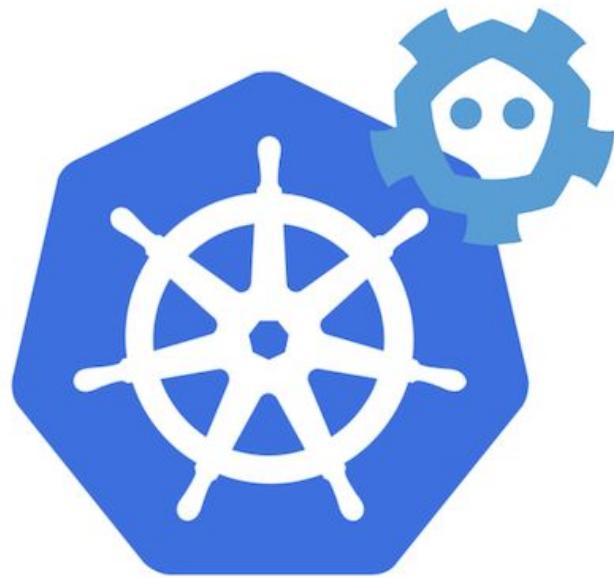
Get Nome

"Clevin"

Tabular - Bancos de Dados Relacionais

Nome	Idade	Localização
Berg	18	Sobral
Bruno	45	Cabrobró
Samuel	32	Salvador





Patroni



Instalação do ETCD

1. Download do binário

```
curl -L https://github.com/etcd-io/etcd/releases/download/v3.3.11/etcd-  
v3.3.11-linux-amd64.tar.gz -o etcd-v3.3.11-linux-amd64.tar.gz
```

2. Extração

```
tar xzvf etcd-v3.3.11-linux-amd64.tar.gz
```

3. Execute o serviço ETCD

```
./etcd
```

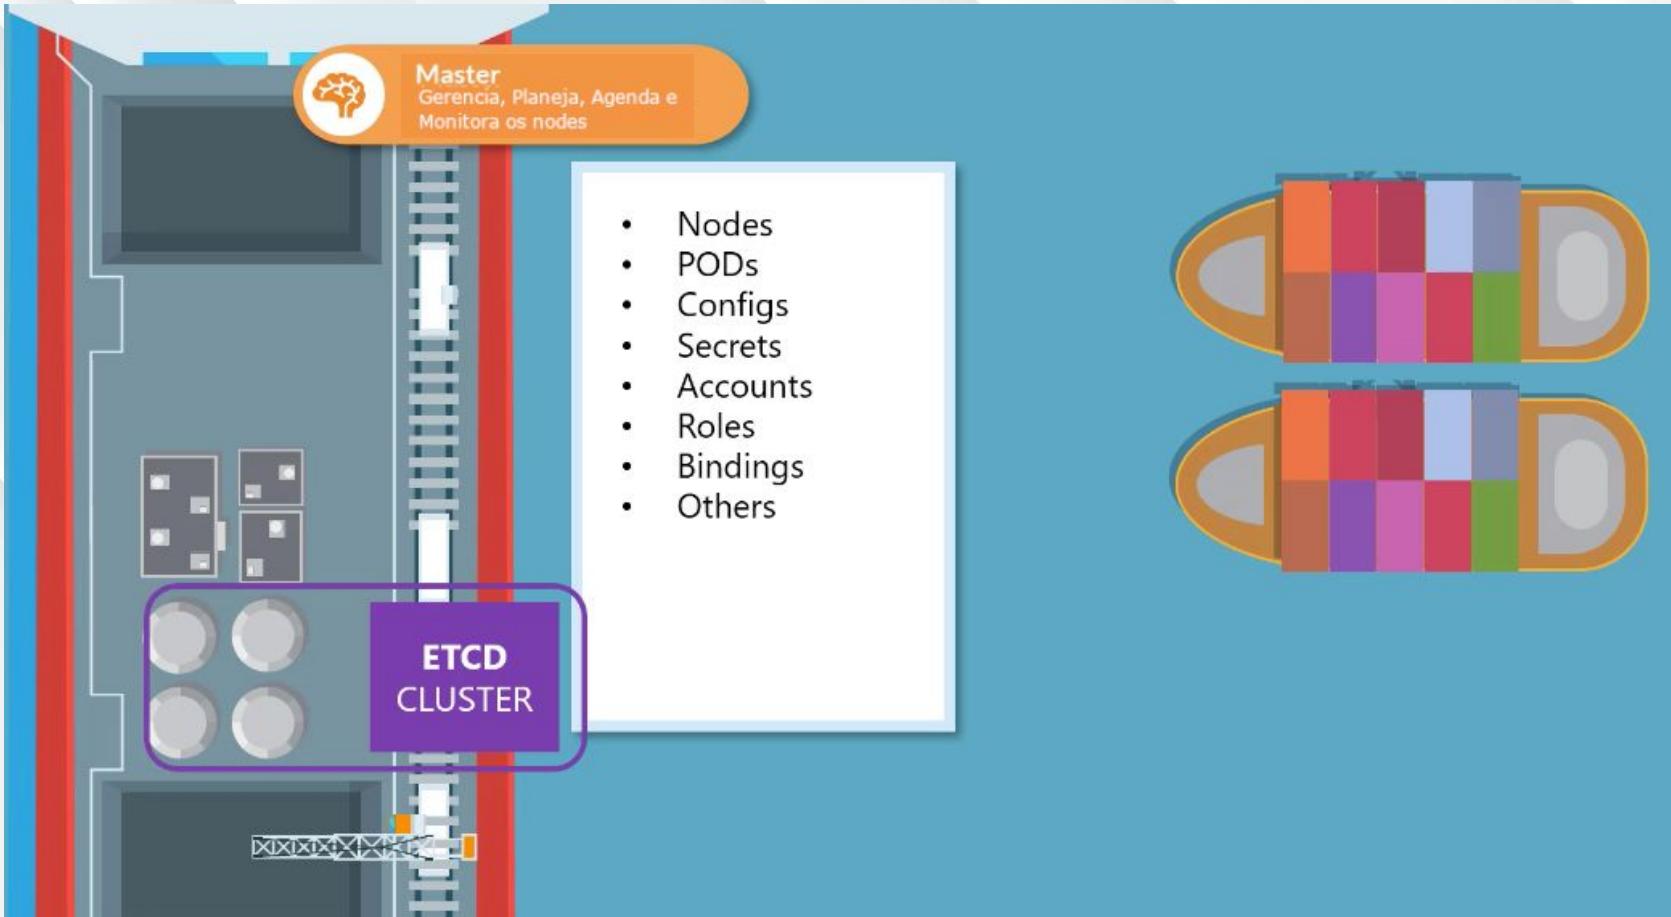
4. Para gerenciar o ETCD, utilizamos o cliente chamado etcdctl:

```
→ etcdctl put chave1 valor1
```

```
→ etcdctl get chave1  
chave1  
valor1
```



ETCD no Kubernetes



Kubeadm

```
kubectl get pods -n kube-system
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
kube-system	coredns-78fcdf6894-prwvl	1/1	Running	0	1h
kube-system	coredns-78fcdf6894-vqd9w	1/1	Running	0	1h
kube-system	etcd-master	1/1	Running	0	1h
kube-system	kube-apiserver-master	1/1	Running	0	1h
kube-system	kube-controller-manager-master	1/1	Running	0	1h
kube-system	kube-proxy-f6k26	1/1	Running	0	1h
kube-system	kube-proxy-hnzsw	1/1	Running	0	1h
kube-system	kube-scheduler-master	1/1	Running	0	1h
kube-system	weave-net-924k8	2/2	Running	1	1h
kube-system	weave-net-hzfcz	2/2	Running	1	1h

```
→ kubectl get pods -n kube-system
```

NAME	READY	STATUS	RESTARTS	AGE
coredns-5644d7b6d9-2kw9p	1/1	Running	25	209d
coredns-5644d7b6d9-gdfqh	1/1	Running	25	209d
etcd-docker-desktop	1/1	Running	48	209d
kube-apiserver-docker-desktop	1/1	Running	47	209d
kube-controller-manager-docker-desktop	1/1	Running	27	209d
kube-proxy-6b9gz	1/1	Running	25	209d
kube-scheduler-docker-desktop	1/1	Running	27	209d
storage-provisioner	1/1	Running	50	209d
vpnkit-controller	1/1	Running	26	209d

Docker

ETCD em alta disponibilidade



```
kubectl exec etcd-master -n kube-system etcdctl get / --prefix -keys-only
/registry/apiregistration.k8s.io/apiservices/v1.
/registry/apiregistration.k8s.io/apiservices/v1.apps
/registry/apiregistration.k8s.io/apiservices/v1.authentication.k8s.io
/registry/apiregistration.k8s.io/apiservices/v1.authorization.k8s.io
/registry/apiregistration.k8s.io/apiservices/v1.autoscaling
/registry/apiregistration.k8s.io/apiservices/v1.batch
/registry/apiregistration.k8s.io/apiservices/v1.networking.k8s.io
/registry/apiregistration.k8s.io/apiservices/v1.rbac.authorization.k8s.io
/registry/apiregistration.k8s.io/apiservices/v1.storage.k8s.io
/registry/apiregistration.k8s.io/apiservices/v1beta1.admissionregistration.k8s.io
```

Registry

minions

pods

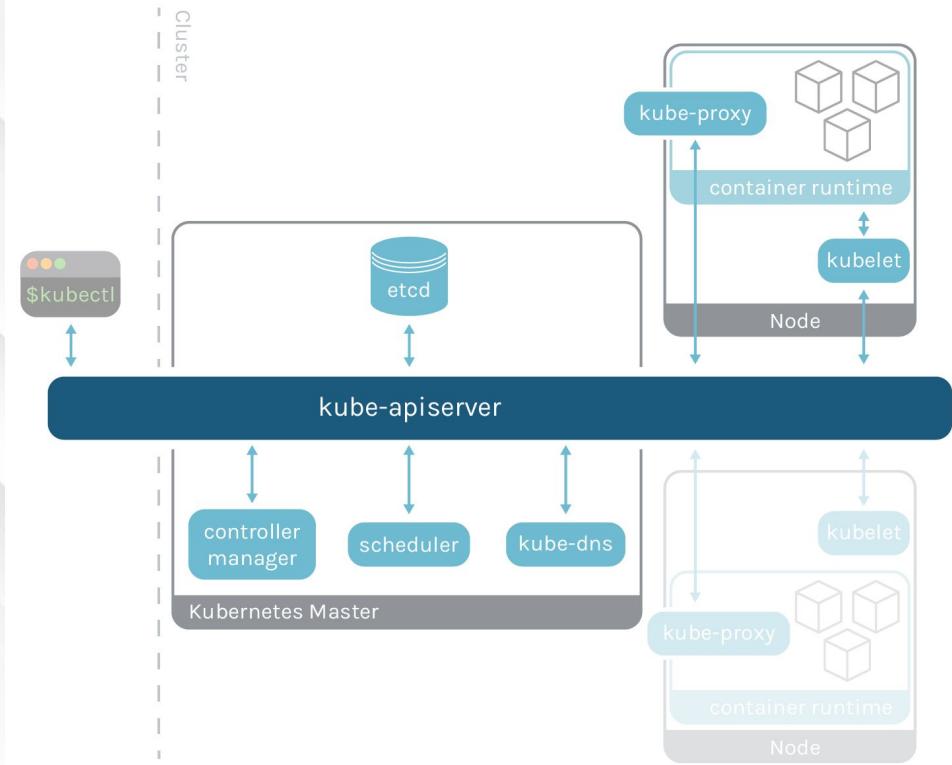
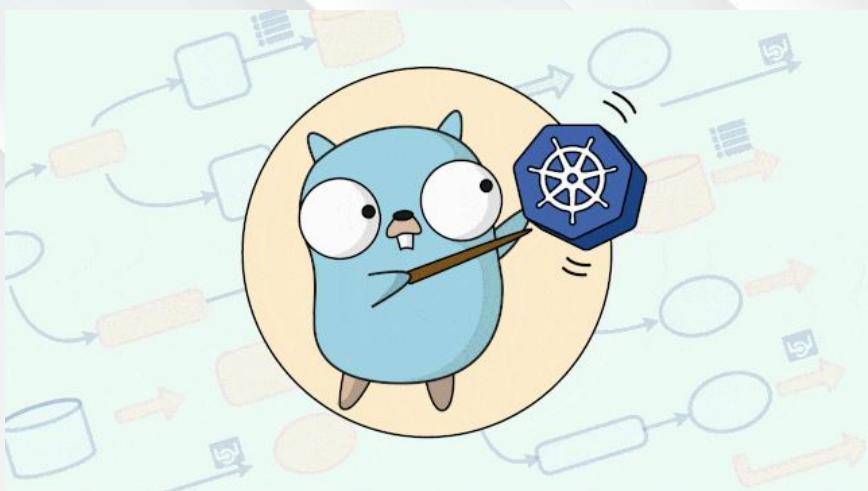
replicasets

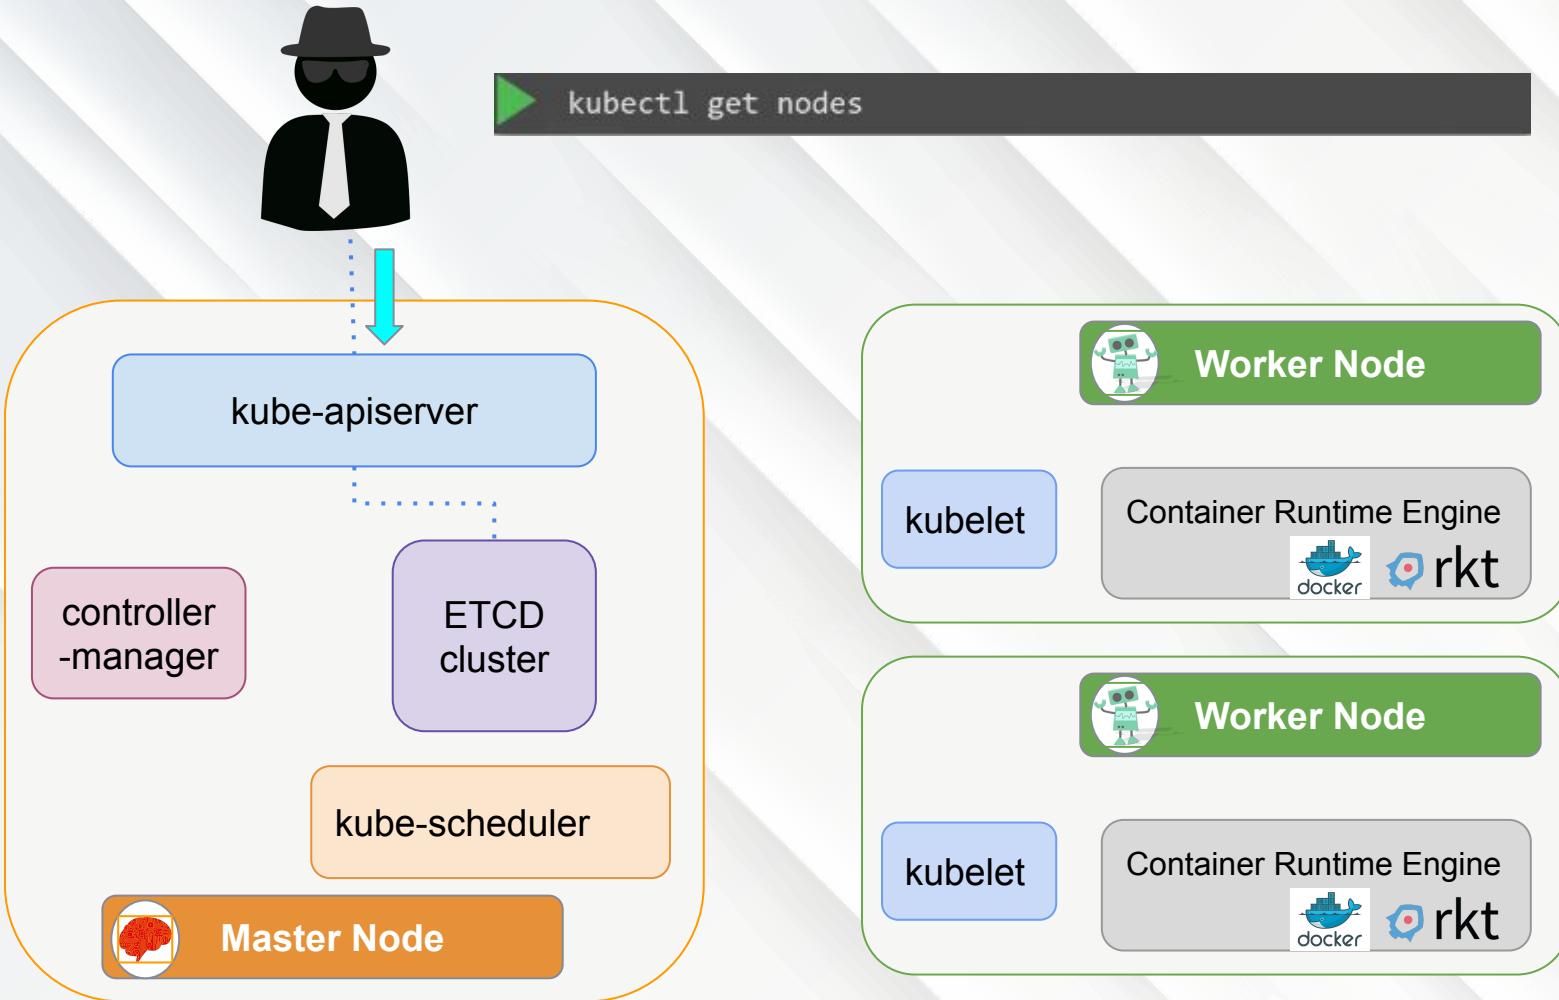
deployments

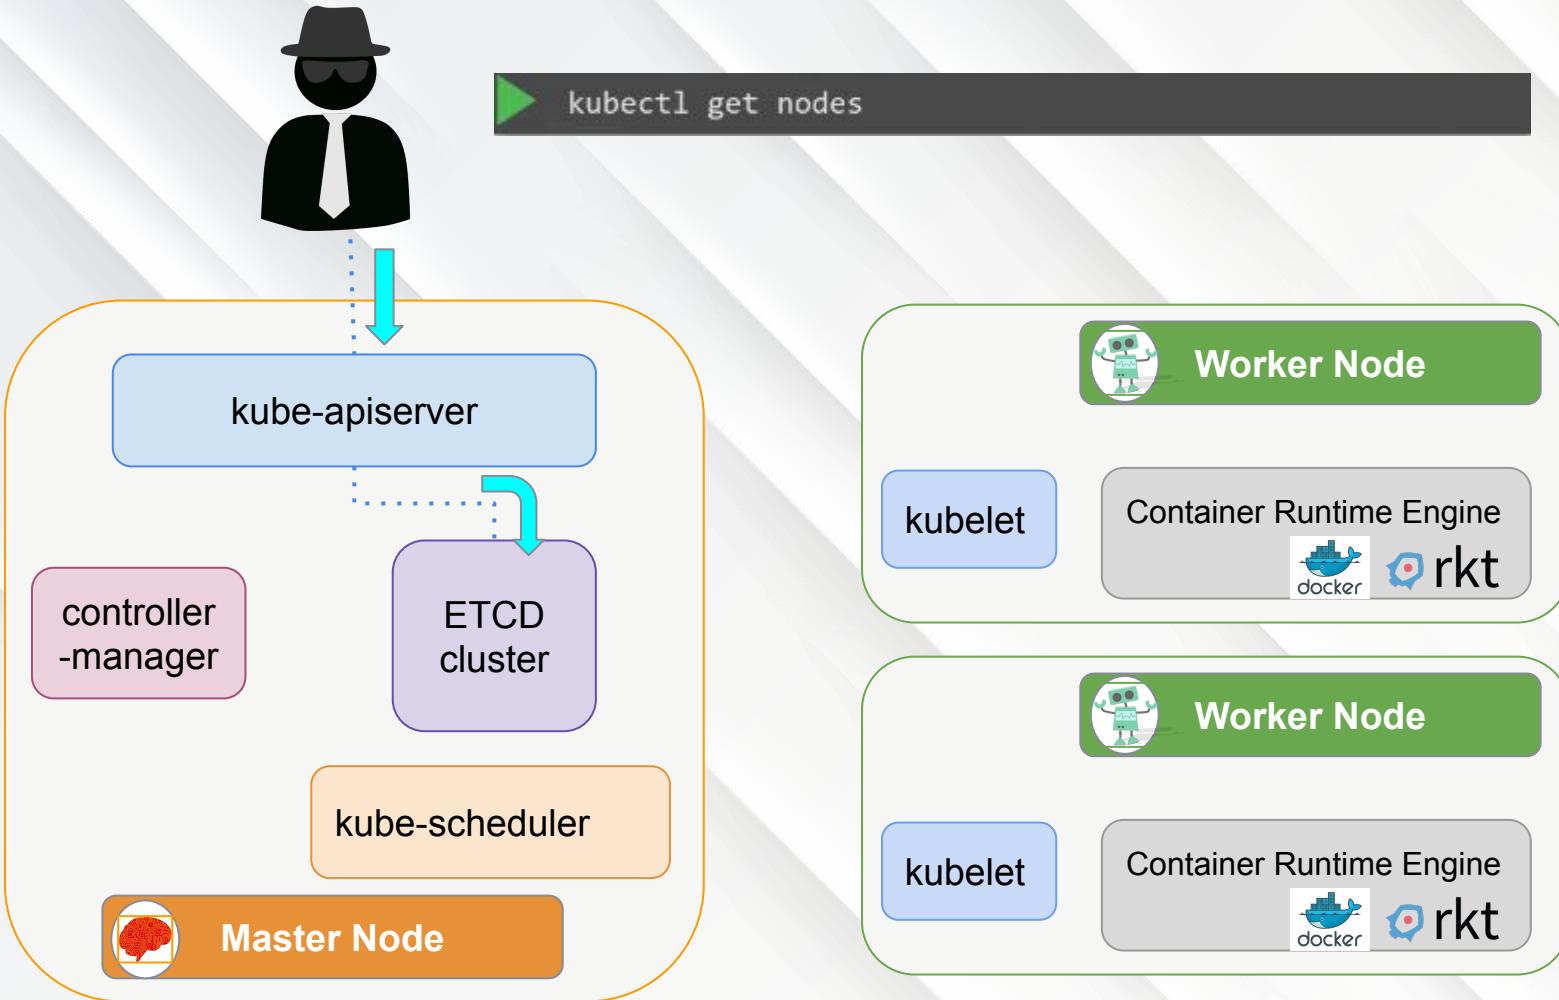
roles

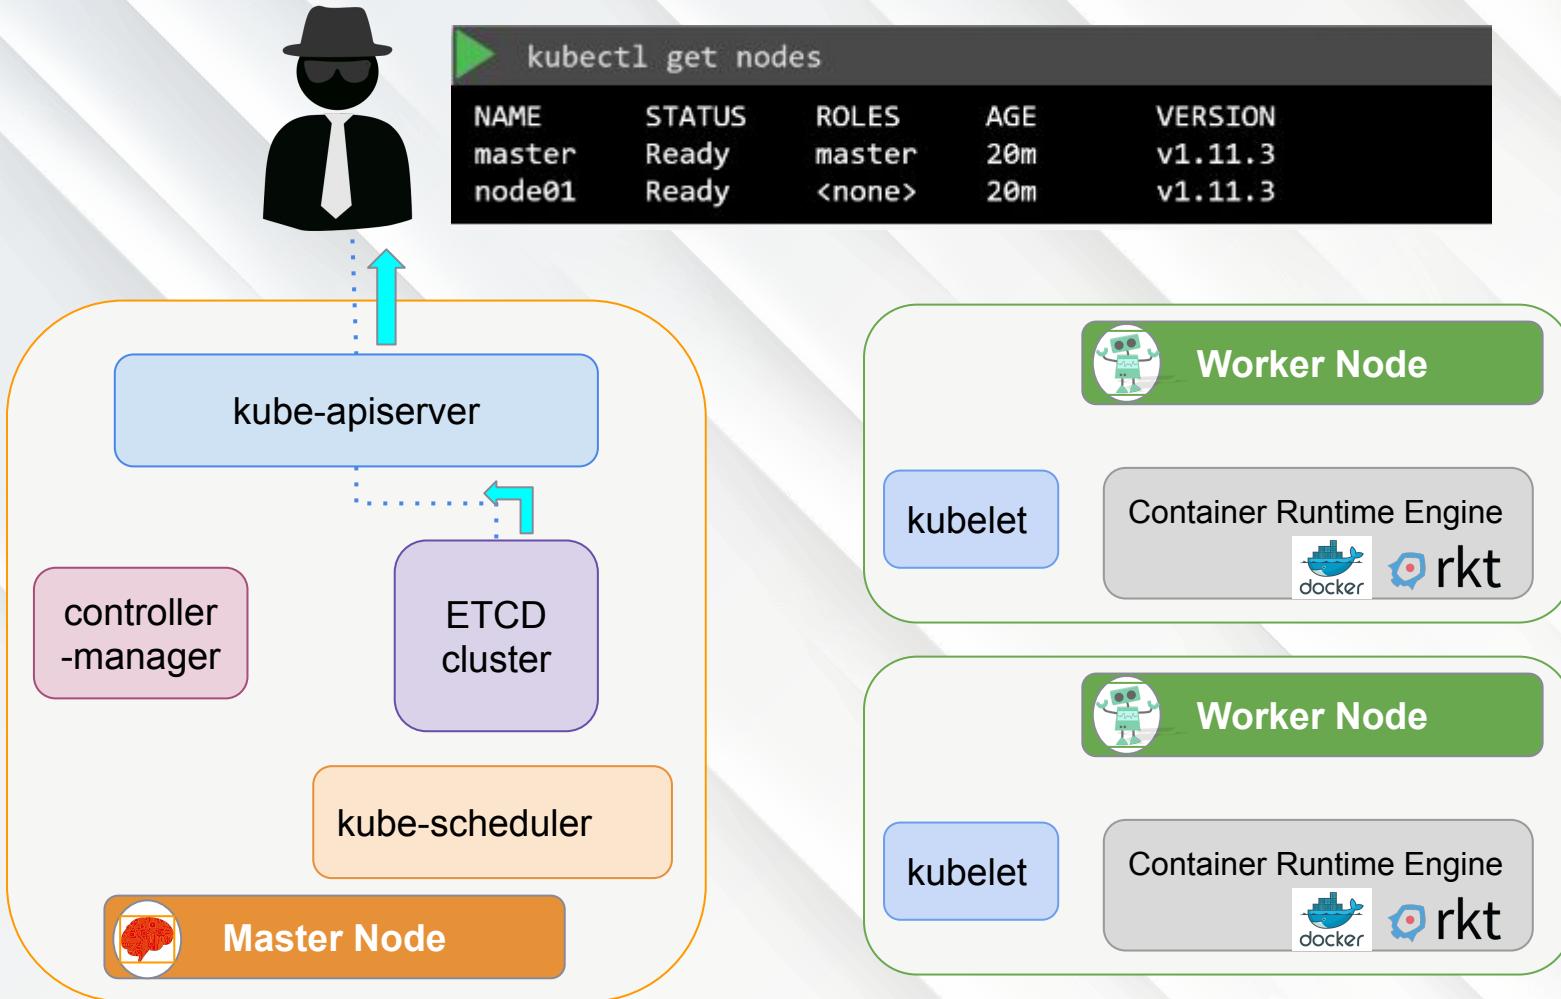
secrets

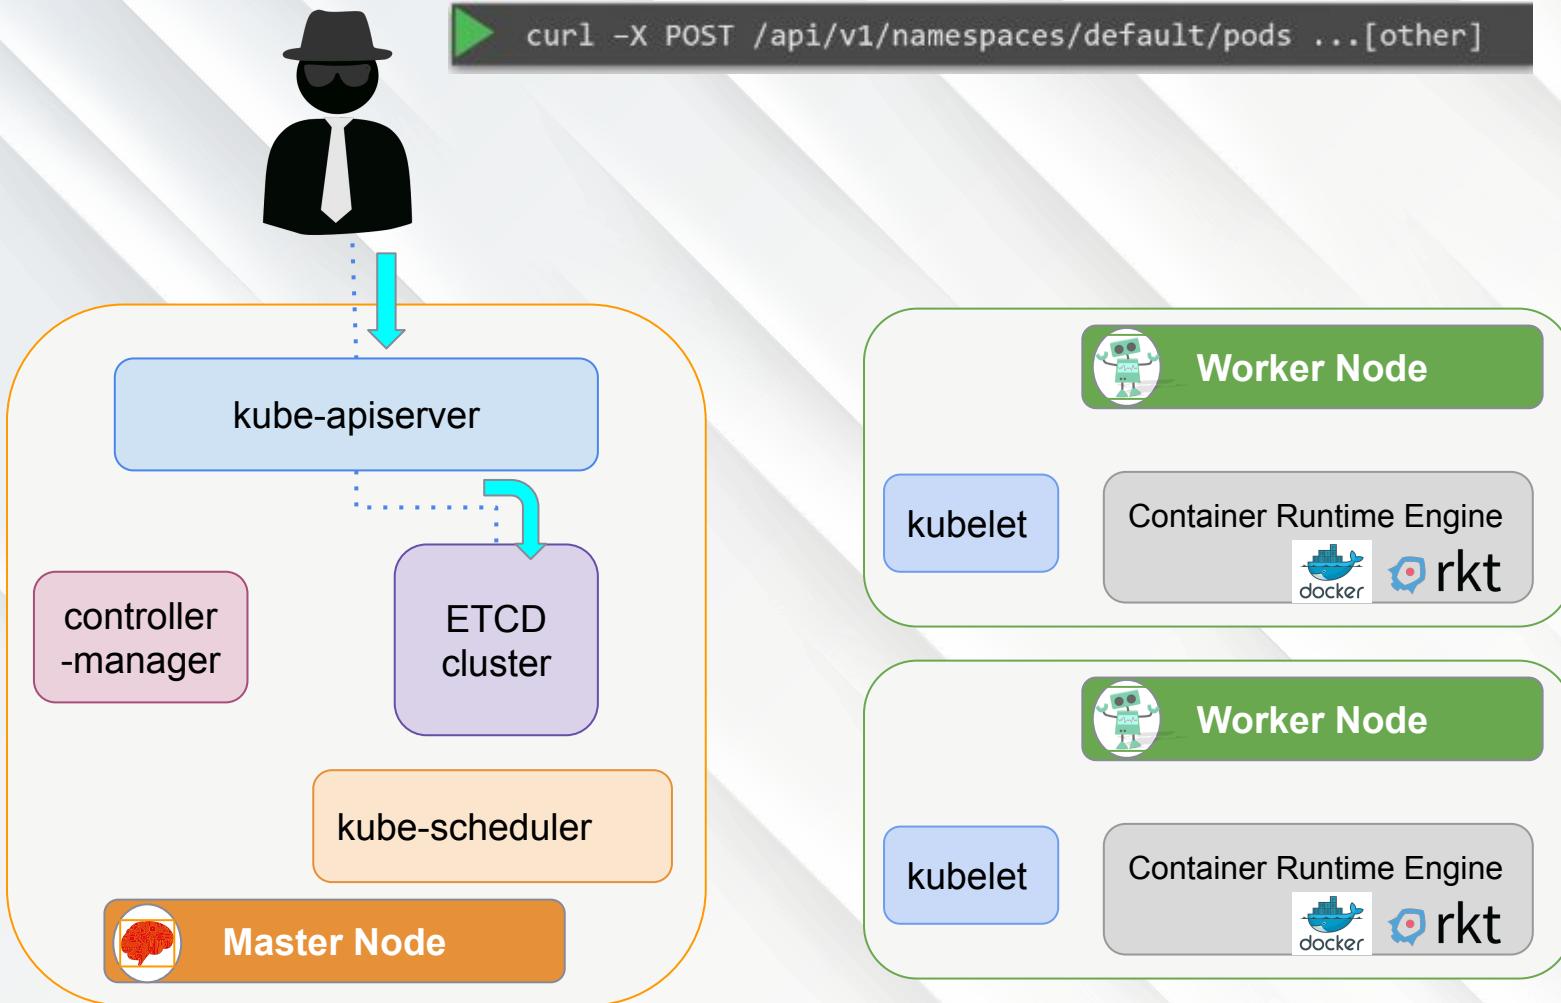
Kube-apiserver

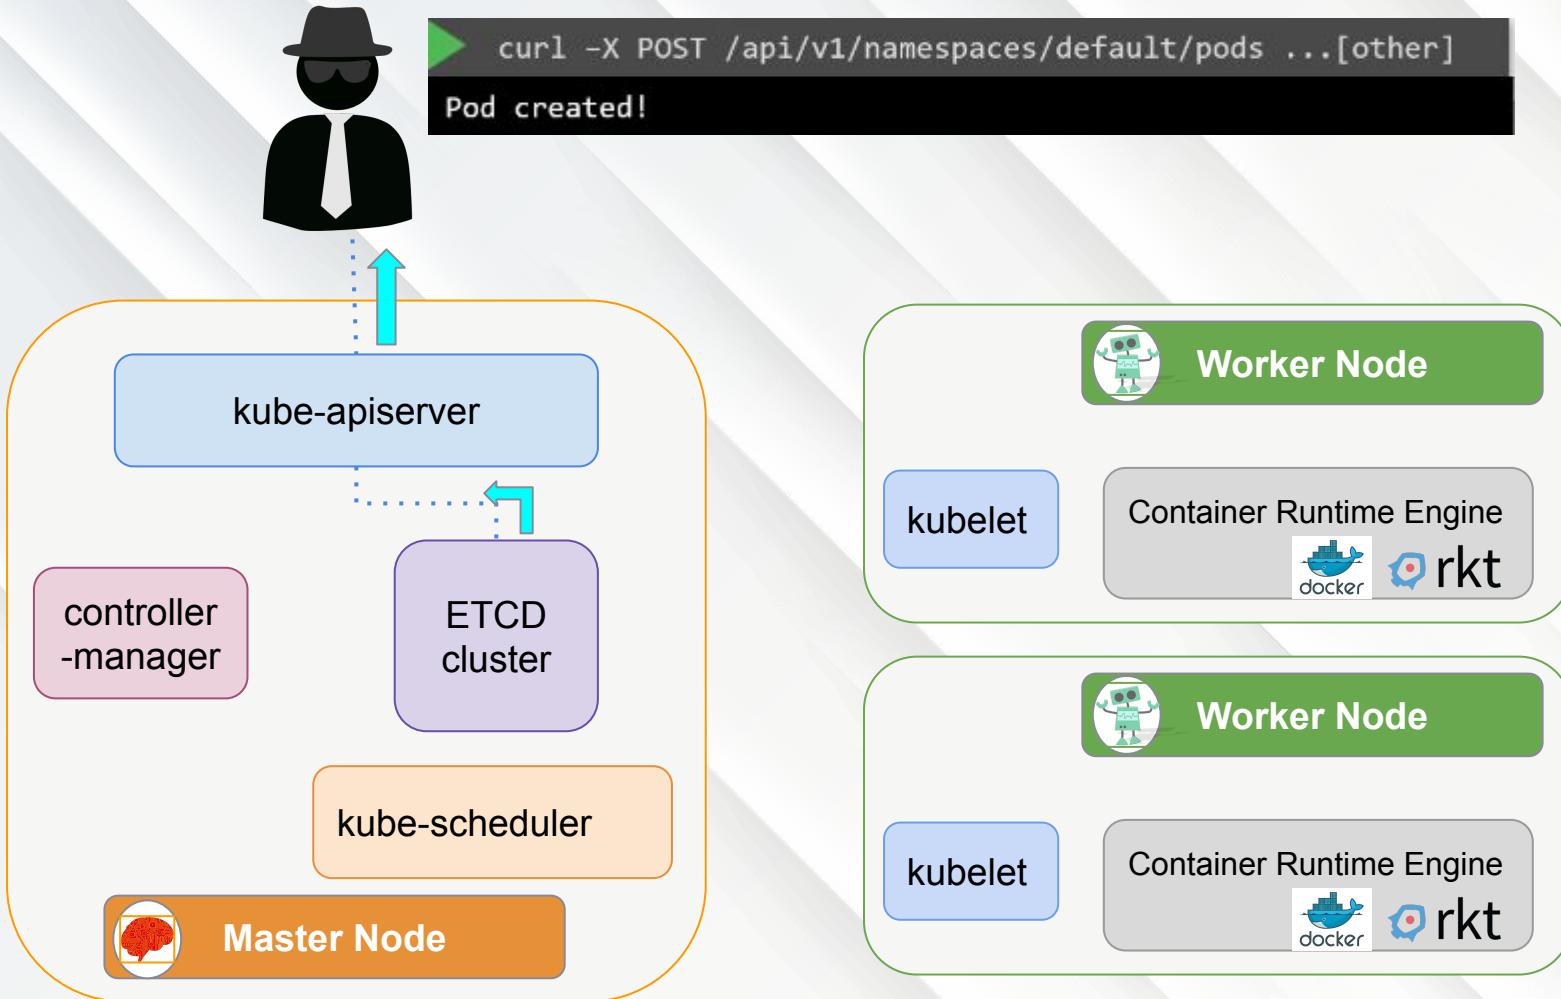








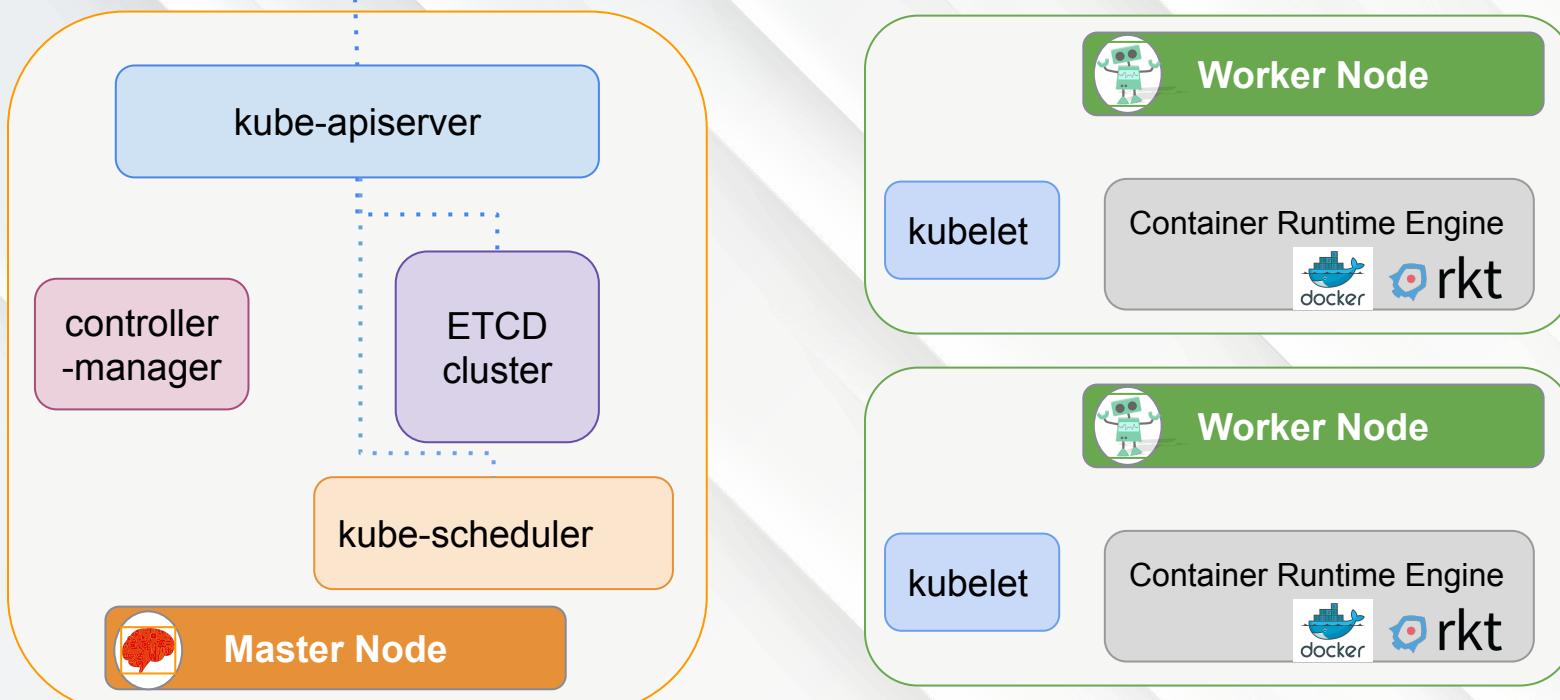






```
curl -X POST /api/v1/namespaces/default/pods ...[other]
```

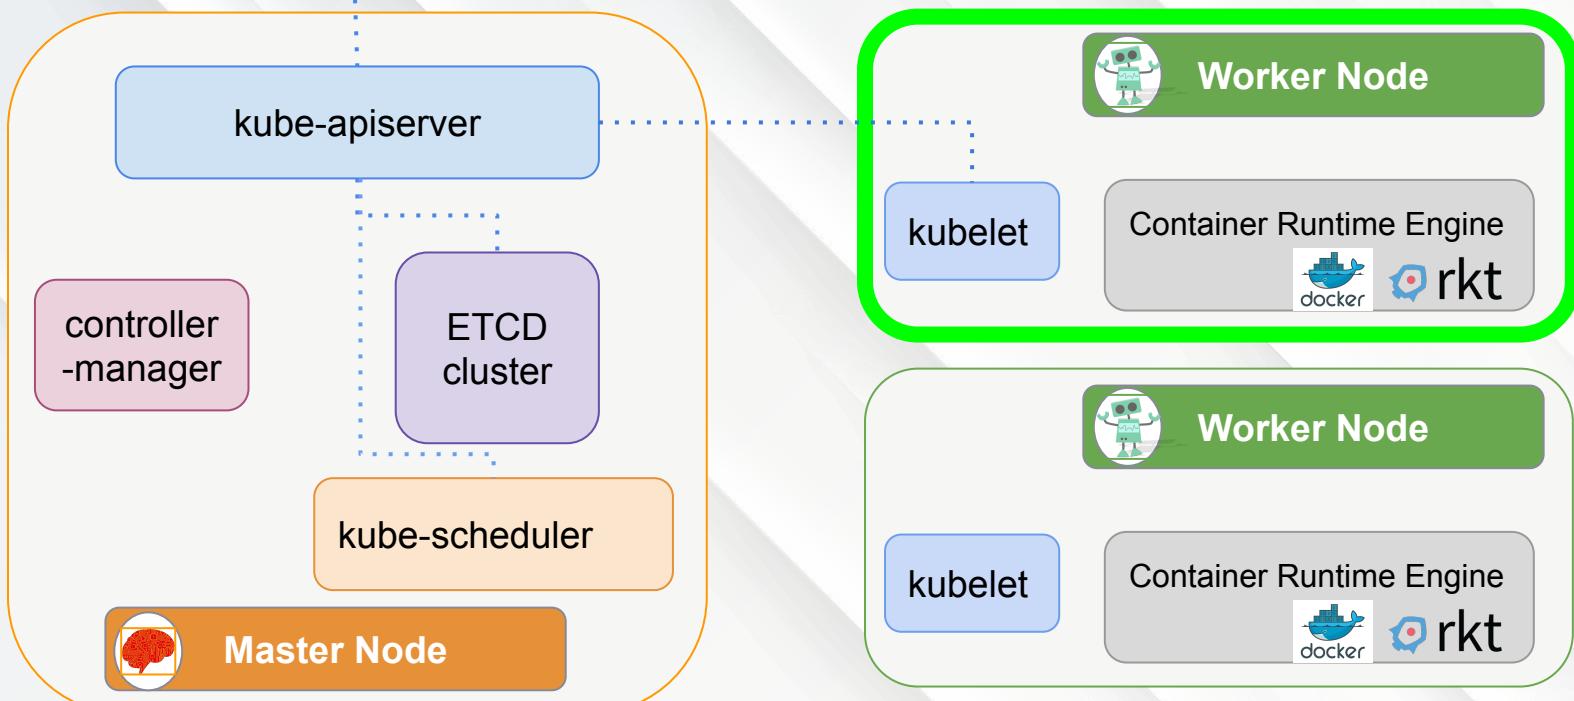
Pod created!





```
curl -X POST /api/v1/namespaces/default/pods ...[other]
```

Pod created!



Kube-apiserver

Autenticar o usuário

Validar as requisições

Recuperar dados (solicitar)

Atualizar o ETCD

Scheduler

Kubelet

Kube-apiserver

```
kubectl get pods -n kube-system
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
kube-system	coredns-78fcdf6894-hwrq9	1/1	Running	0	16m
kube-system	coredns-78fcdf6894-rzhjr	1/1	Running	0	16m
kube-system	etcd-master	1/1	Running	0	15m
kube-system	kube-apiserver-master	1/1	Running	0	15m
kube-system	kube-controller-manager-master	1/1	Running	0	15m
kube-system	kube-proxy-lzt6f	1/1	Running	0	16m
kube-system	kube-proxy-zm5qd	1/1	Running	0	16m
kube-system	kube-scheduler-master	1/1	Running	0	15m
kube-system	weave-net-29z42	2/2	Running	1	16m
kube-system	weave-net-snmdl				

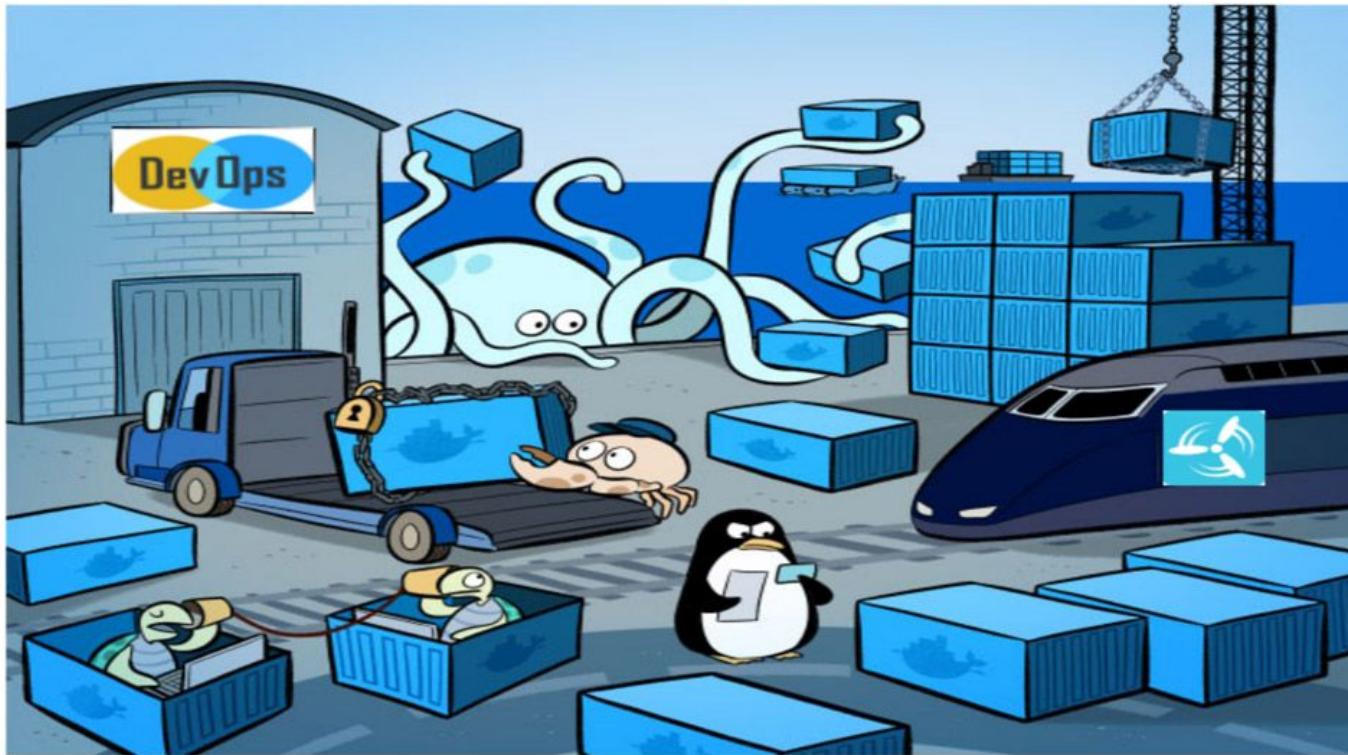
Kubeadm

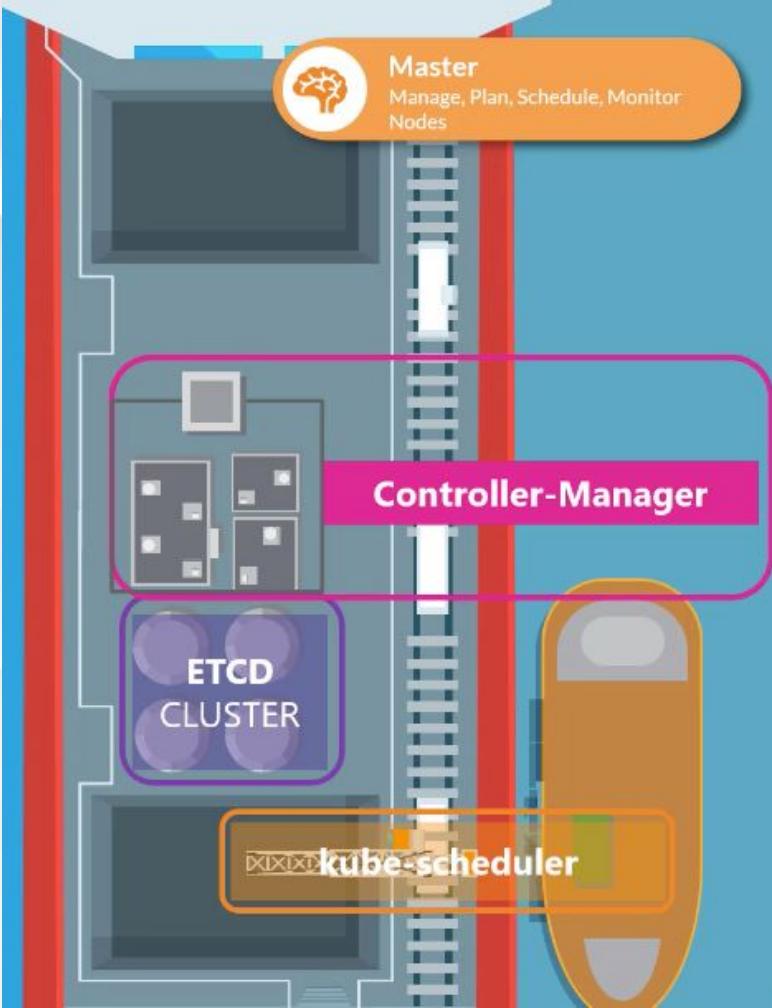
```
+ kubectl get pods -n kube-system
```

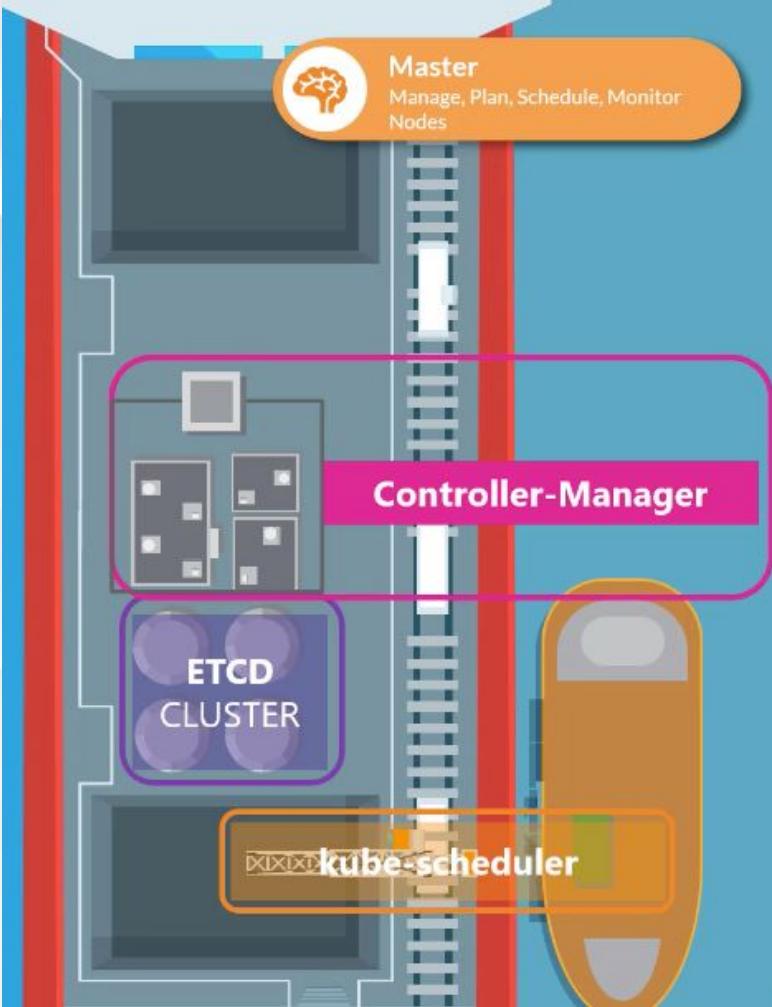
NAME	READY	STATUS	RESTARTS	AGE
coredns-5644d7b6d9-2kw9p	1/1	Running	26	209d
coredns-5644d7b6d9-gdfqh	1/1	Running	26	209d
etcd-docker-desktop	1/1	Running	49	209d
kube-apiserver-docker-desktop	1/1	Running	48	209d
kube-controller-manager-docker-desktop	1/1	Running	28	209d
kube-proxy-6b9gz	1/1	Running	26	209d
kube-scheduler-docker-desktop	1/1	Running	28	209d
storage-provisioner	1/1	Running	52	209d
vpnkit-controller	1/1	Running	27	209d

Docker

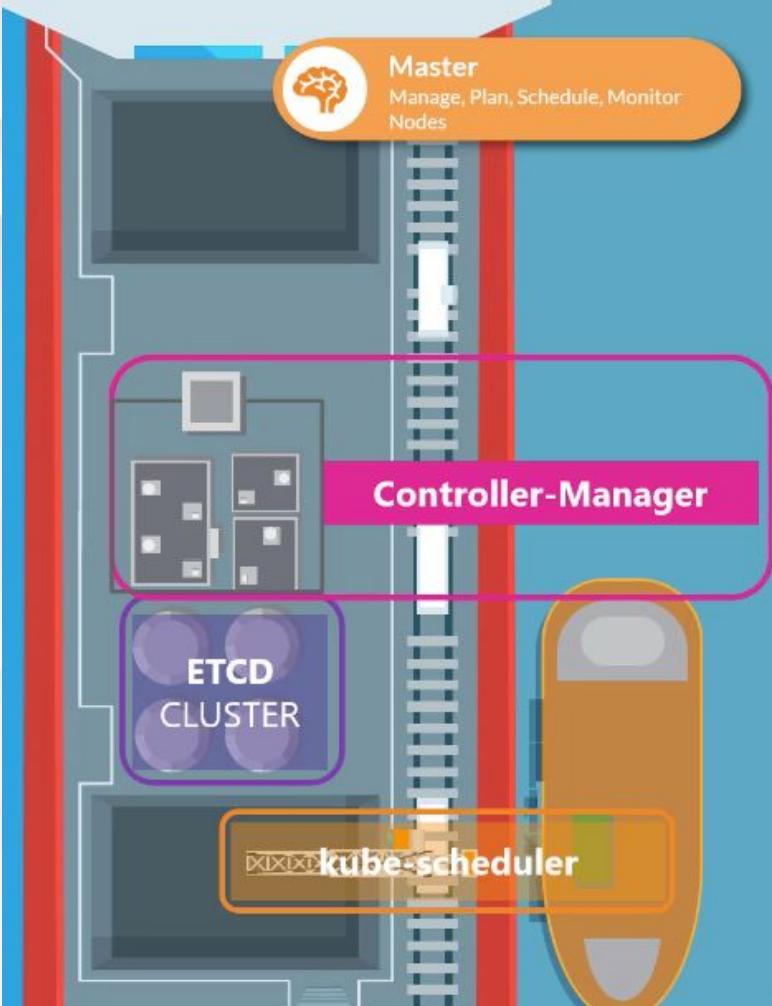
Kube Controller Manager





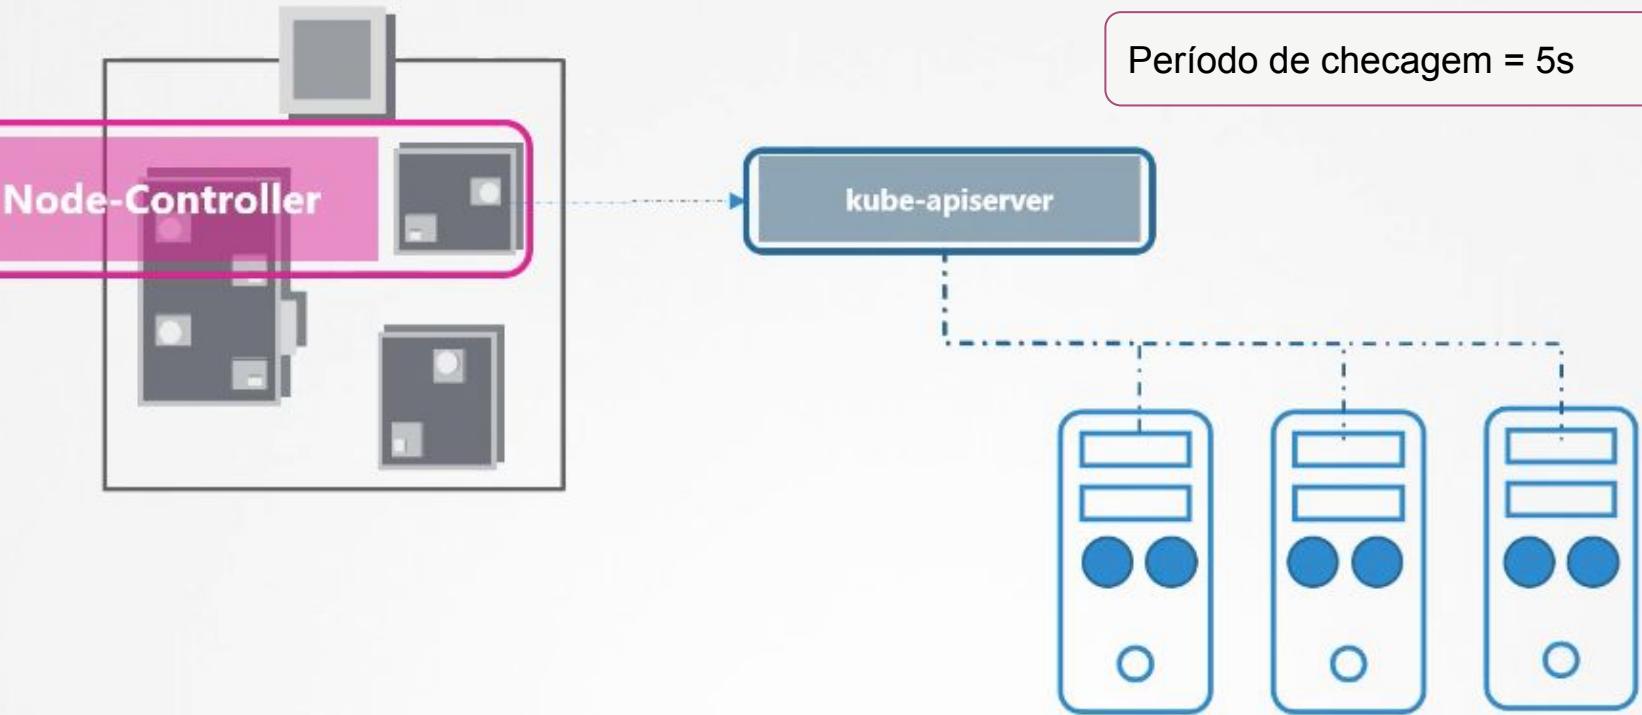


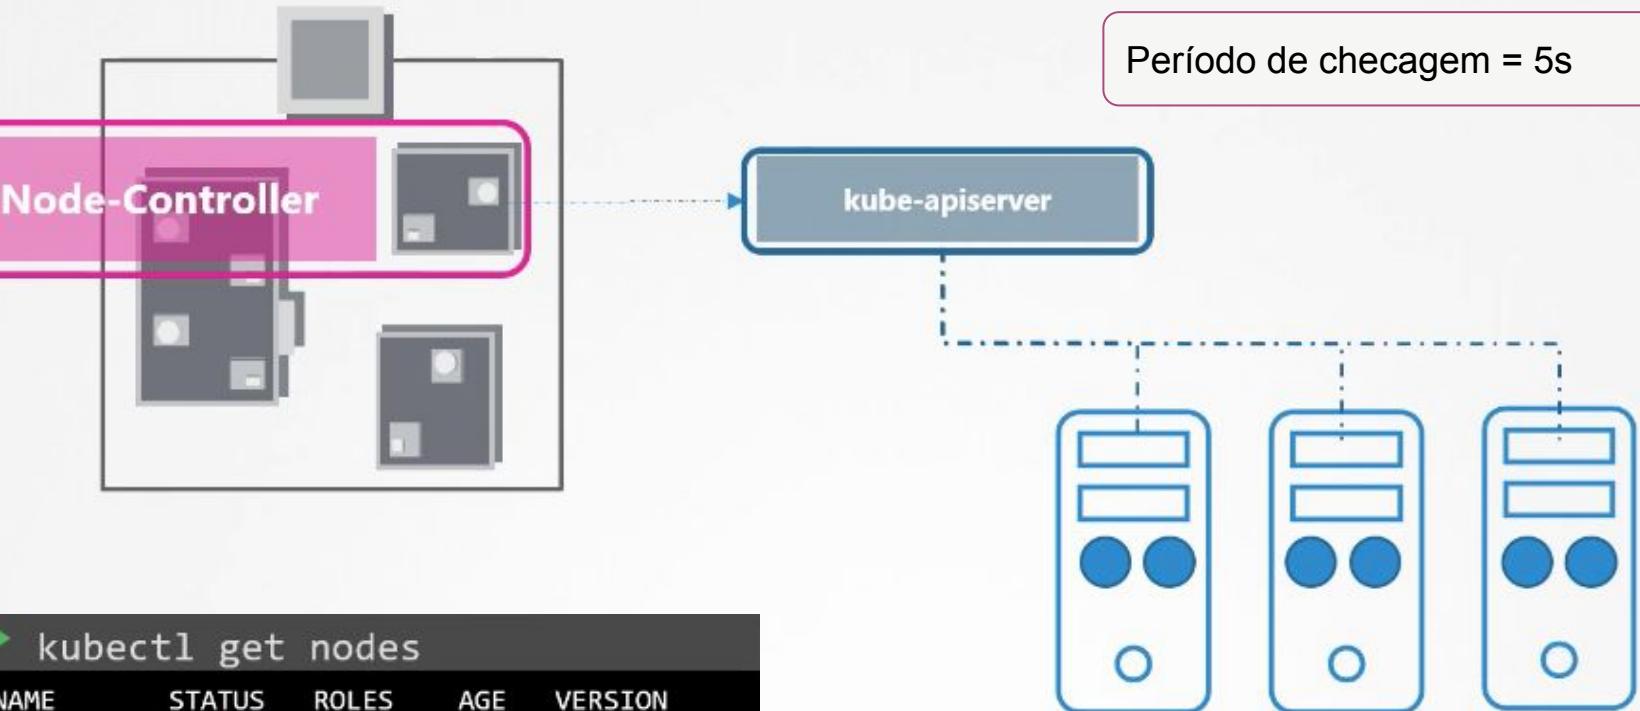
Verificando Status



Verificando Status

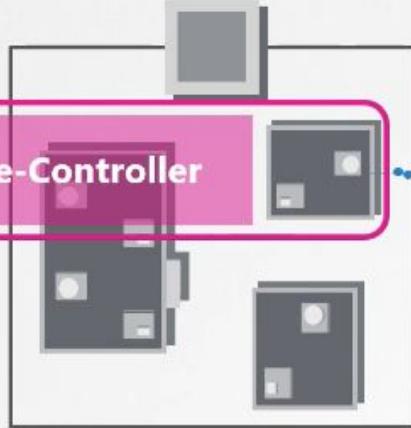
Remediando situações





```
▶ kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
worker-1	Ready	<none>	8d	v1.13.0
worker-2	Ready	<none>	8d	v1.13.0



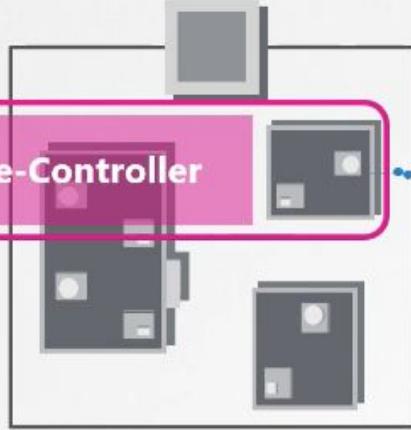
Período de checagem = 5s

Grace Period = 40s

POD Eviction Timeout = 5m

```
kubectl get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
worker-1	Ready	<none>	8d	v1.13.0
worker-2	NotReady	<none>	8d	v1.13.0

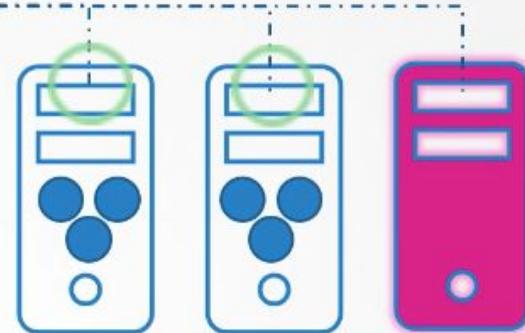


kube-apiserver

Período de checagem = 5s

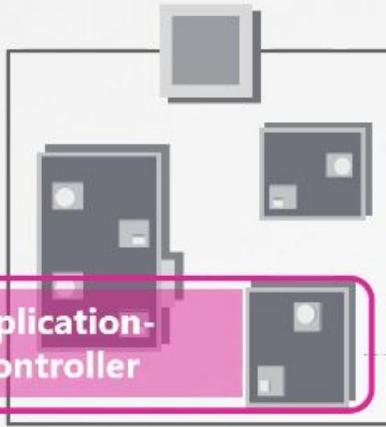
Grace Period = 40s

POD Eviction Timeout = 5m



▶ kubectl get nodes

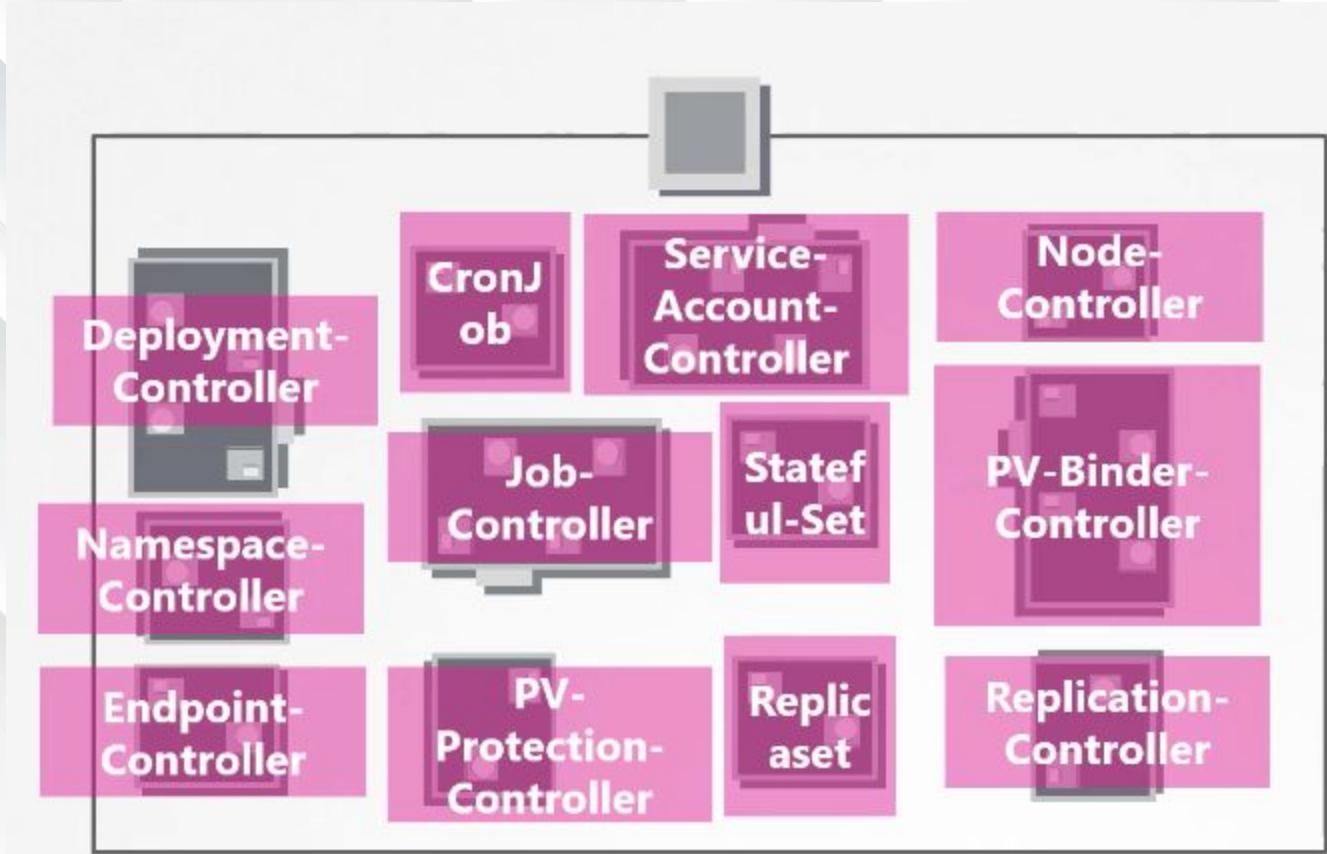
NAME	STATUS	ROLES	AGE	VERSION
worker-1	Ready	<none>	8d	v1.13.0
worker-2	NotReady	<none>	8d	v1.13.0



Período de checagem = 5s

Grace Period = 40s

POD Eviction Timeout = 5m



Kube-Controller-Manager

```
kubectl get pods -n kube-system
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
kube-system	coredns-78fcdf6894-hwrq9	1/1	Running	0	16m
kube-system	coredns-78fcdf6894-rzhjr	1/1	Running	0	16m
kube-system	etcd-master	1/1	Running	0	15m
kube-system	kube-apiserver-master	1/1	Running	0	15m
kube-system	kube-controller-manager-master	1/1	Running	0	15m
kube-system	kube-proxy-lzt6f	1/1	Running	0	16m
kube-system	kube-proxy-zm5qd	1/1	Running	0	16m
kube-system	kube-scheduler-master	1/1	Running	0	15m
kube-system	weave-net-29z42	2/2	Running	1	16m
kube-system	weave-net-snmdl	2/2	Running	1	16m

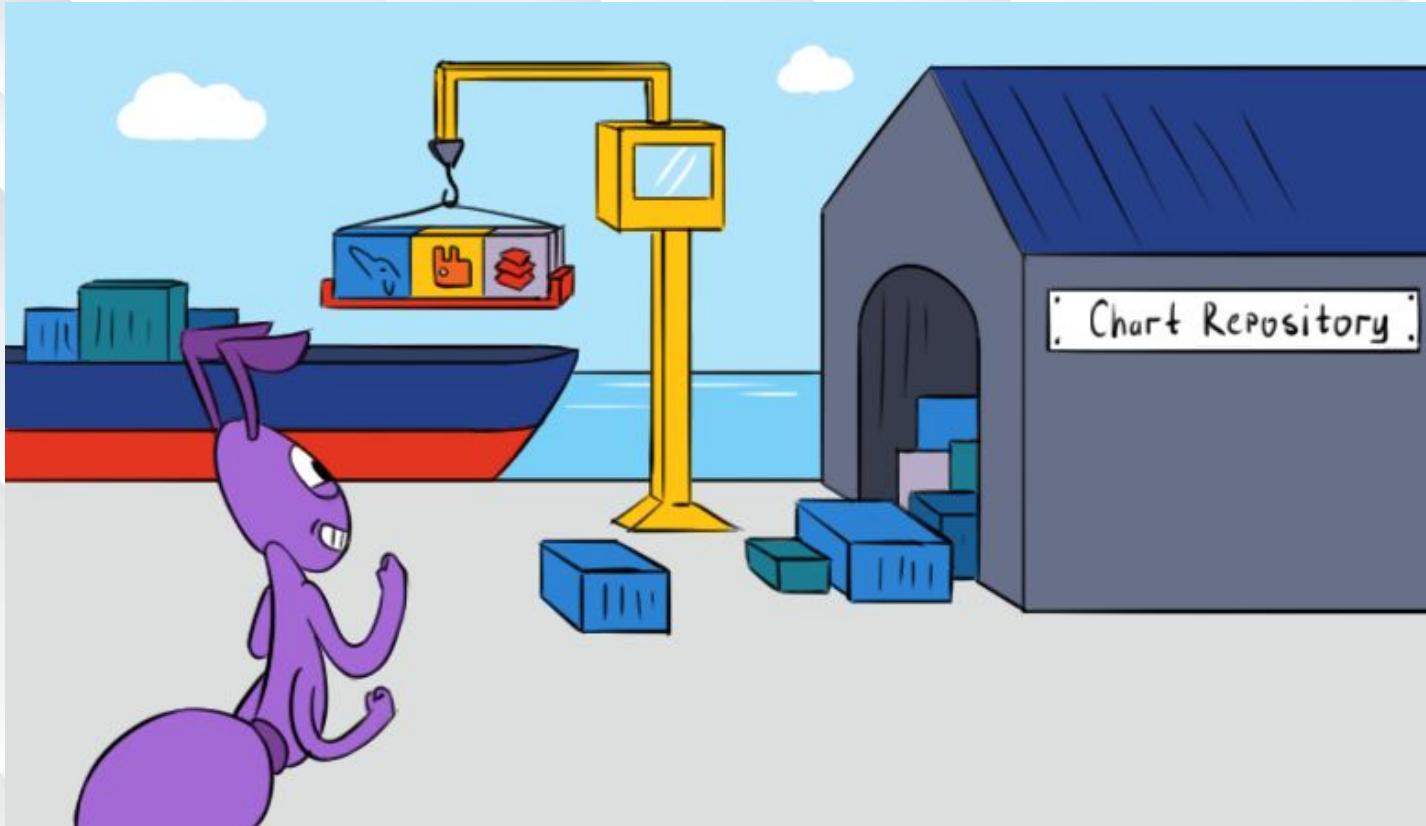
Kubeadm

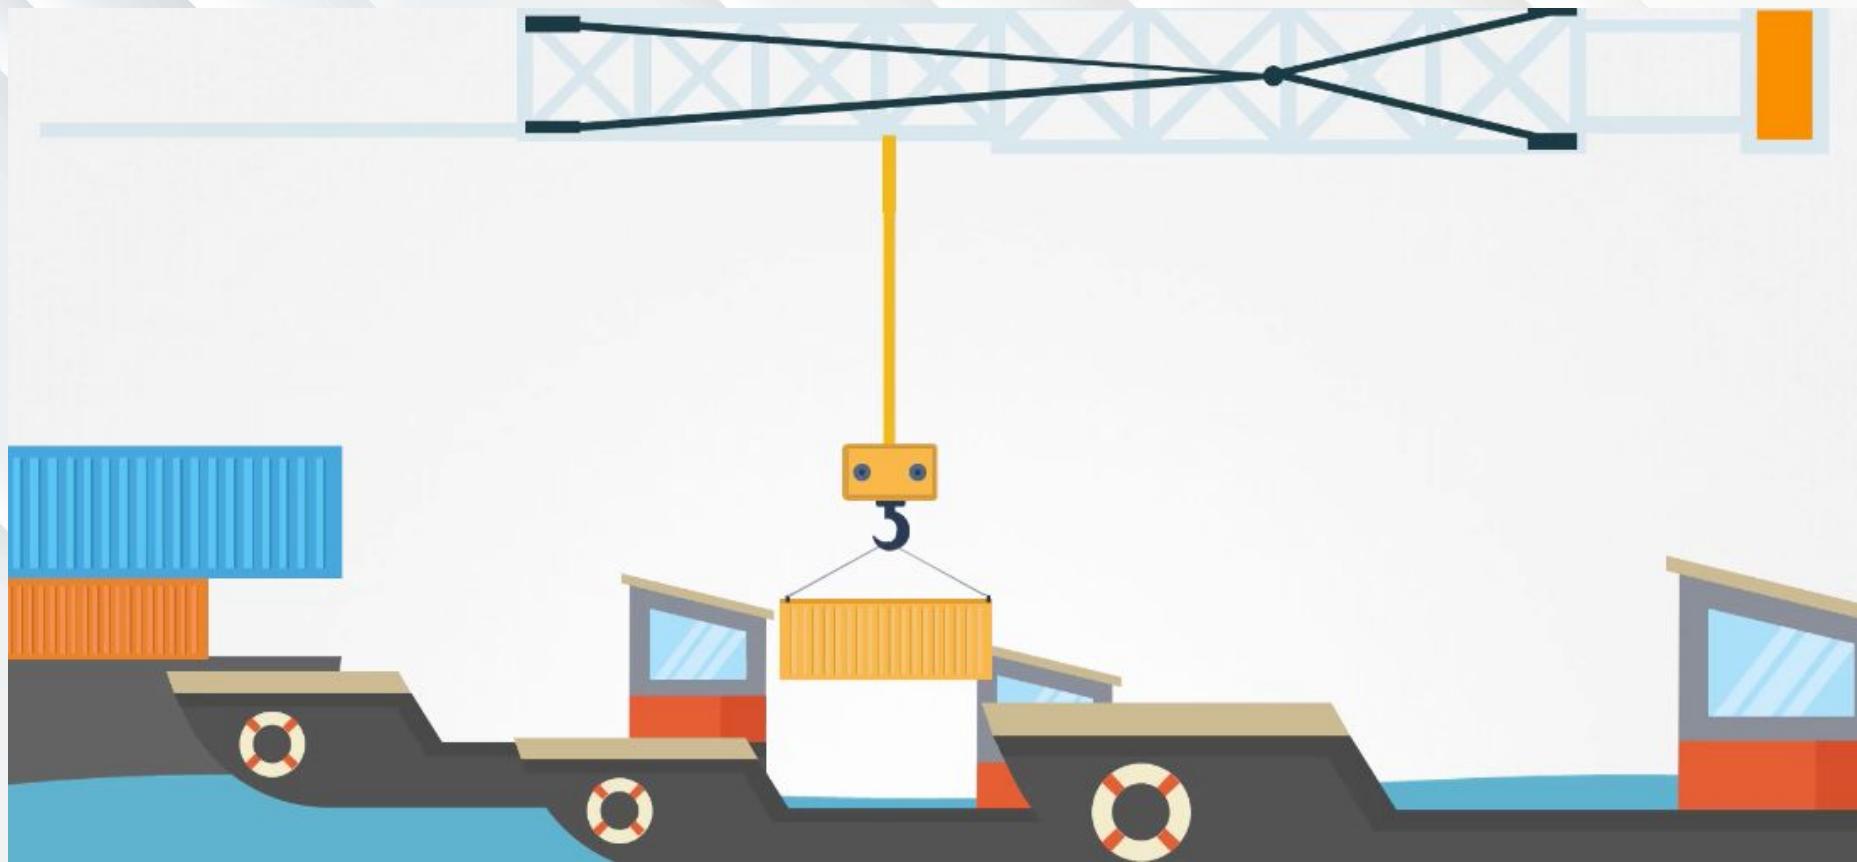
```
+ kubectl get pods -n kube-system
```

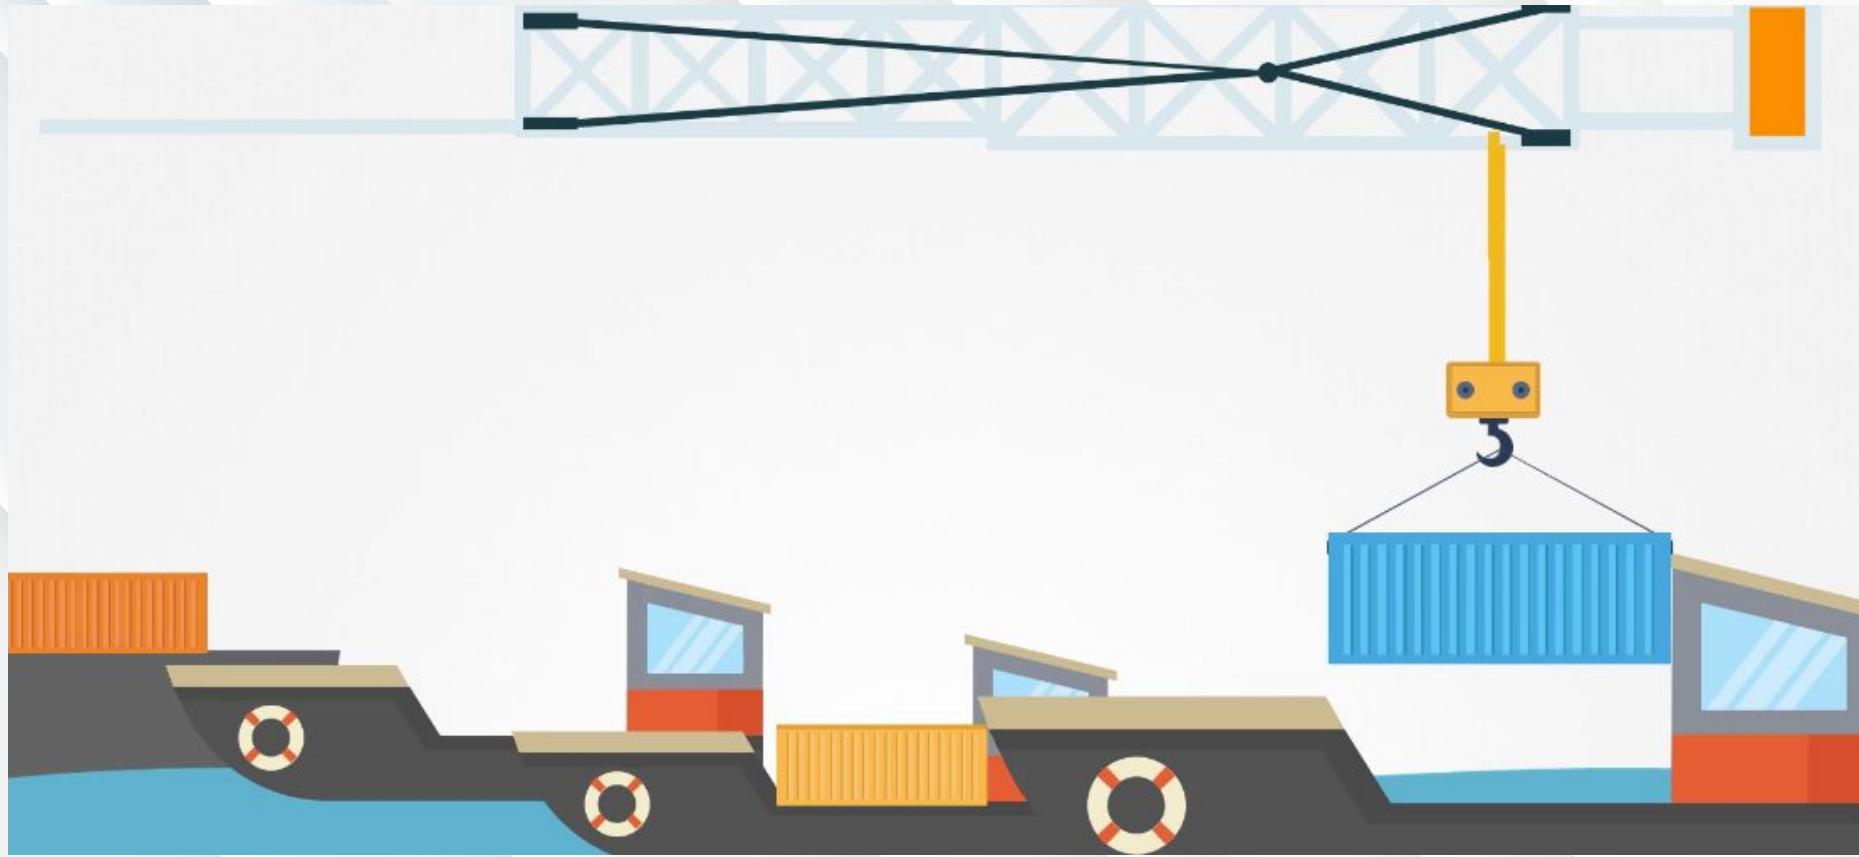
NAME	READY	STATUS	RESTARTS	AGE
coredns-5644d7b6d9-2kw9p	1/1	Running	26	209d
coredns-5644d7b6d9-gdfqh	1/1	Running	26	209d
etcd-docker-desktop	1/1	Running	49	209d
kube-apiserver-docker-desktop	1/1	Running	48	209d
kube-controller-manager-docker-desktop	1/1	Running	28	209d
kube-proxy-6b9gz	1/1	Running	26	209d
kube-scheduler-docker-desktop	1/1	Running	28	209d
storage-provisioner	1/1	Running	52	209d
vpnkit-controller	1/1	Running	27	209d

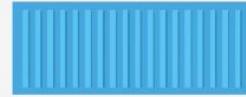
Docker

Kube Scheduler







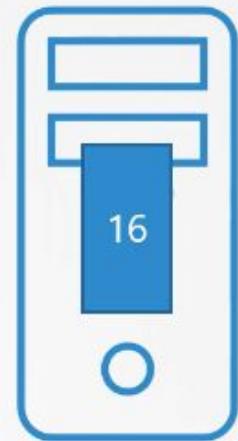
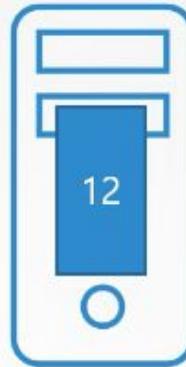
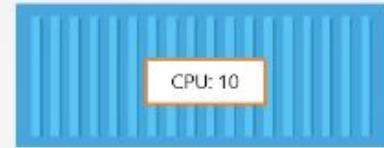


Scheduler verifica o tempo todo os pods que precisam ser alocados e os nodes existentes.

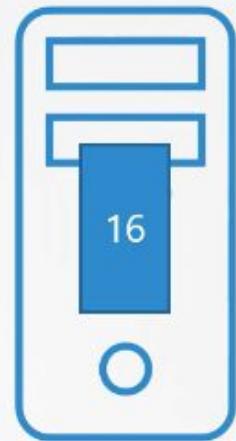
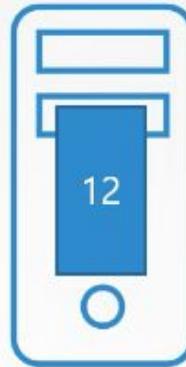
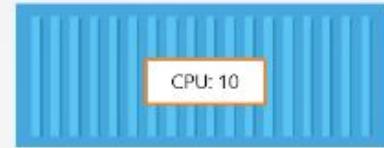
Compara os requisitos e toma as decisões adequadas.



Fase 1: Filtrar os
nodes

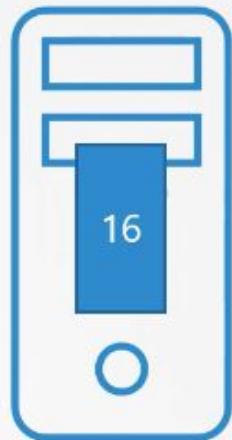
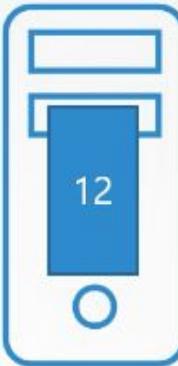
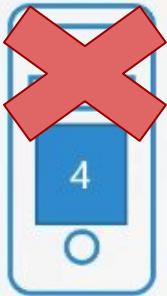
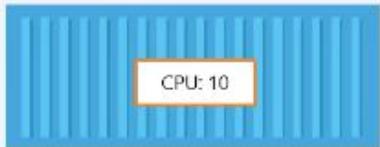


Fase 1: Filtrar os
nodes



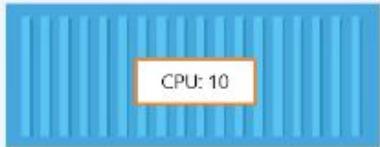
Fase 1: Filtrar os
nodes

Fase 2: Ranking dos
nodes

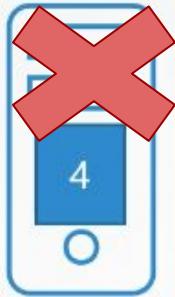


Fase 1: Filtrar os nodes

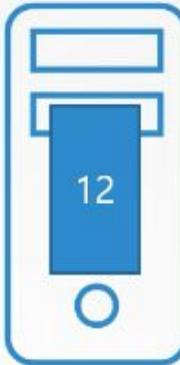
Fase 2: Ranking dos nodes



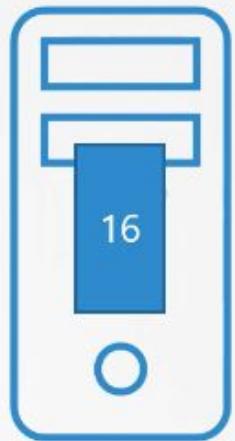
CPU: 10



$$12 - 10 = 2$$

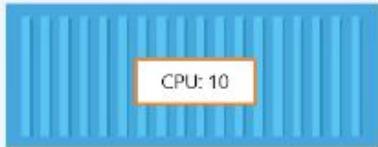


$$16 - 10 = 6$$

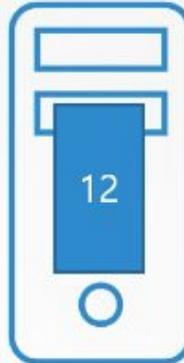


Fase 1: Filtrar os nodes

Fase 2: Ranking dos nodes



$$12 - 10 = 2$$



Adiante:

- Taints
- Tolerations
- Node Selectors
- Affinity

```
→ kubectl get pods -n kube-system
```

NAME	READY	STATUS	RESTARTS	AGE
coredns-5644d7b6d9-2kw9p	1/1	Running	26	209d
coredns-5644d7b6d9-gdfqh	1/1	Running	26	209d
etcd-docker-desktop	1/1	Running	49	209d
kube-apiserver-docker-desktop	1/1	Running	48	209d
kube-controller-manager-docker-desktop	1/1	Running	28	209d
kube-proxy-6b9gz	1/1	Running	26	209d
kube-scheduler-docker-desktop	1/1	Running	28	209d
storage-provisioner	1/1	Running	52	209d
vpnkit-controller	1/1	Running	27	209d

Docker

Kubelet





Master

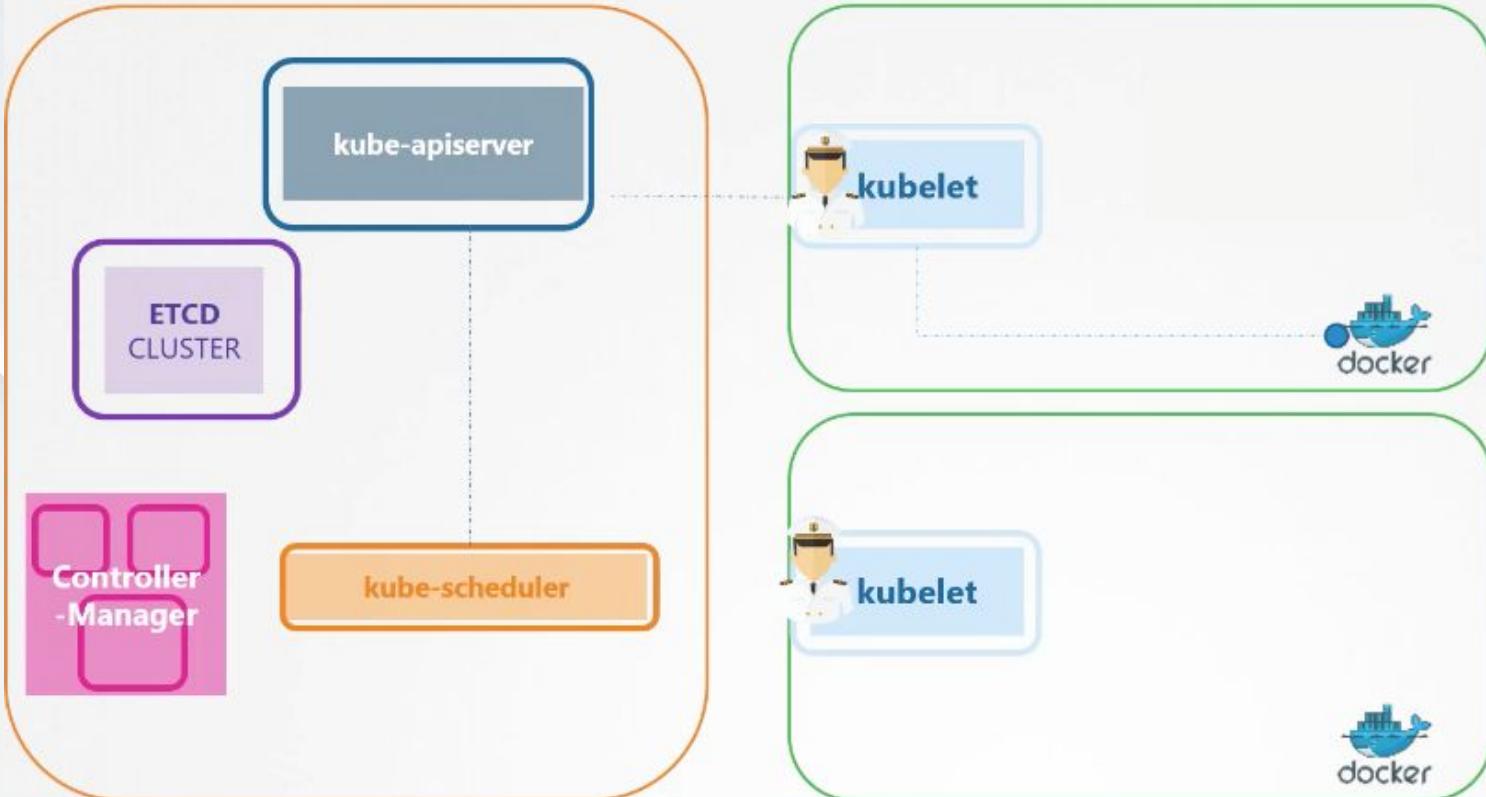
Manage, Plan, Schedule, Monitor
Nodes



Worker Nodes

Host Application as Containers

Registrar um
node





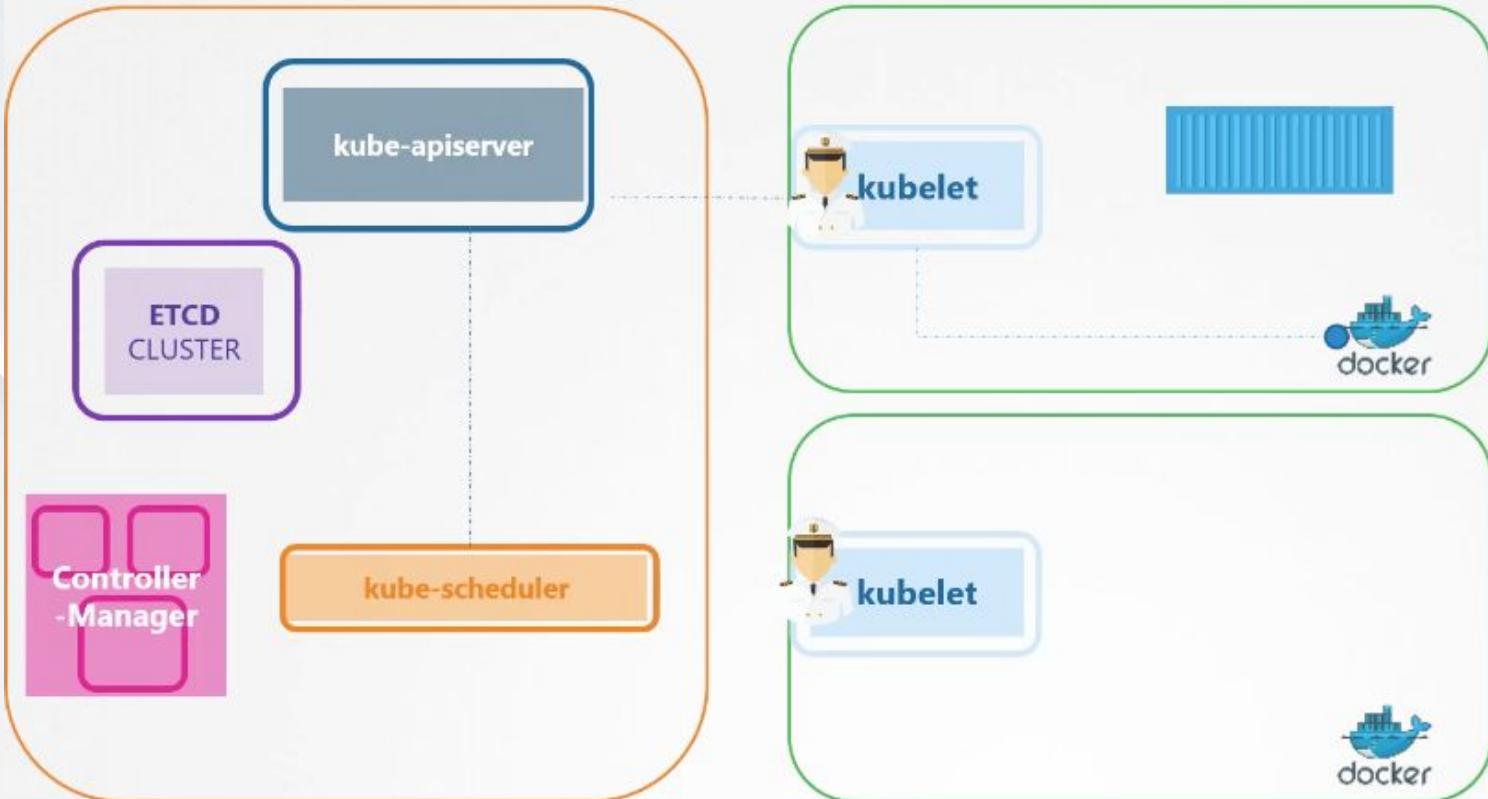
Master

Manage, Plan, Schedule, Monitor
Nodes



Worker Nodes

Host Application as Containers



Registrar um
node

Criar PODs

Instalação do Kubelet

```
wget https://storage.googleapis.com/kubernetes-release/release/v1.13.0/bin/linux/amd64/kubelet
```

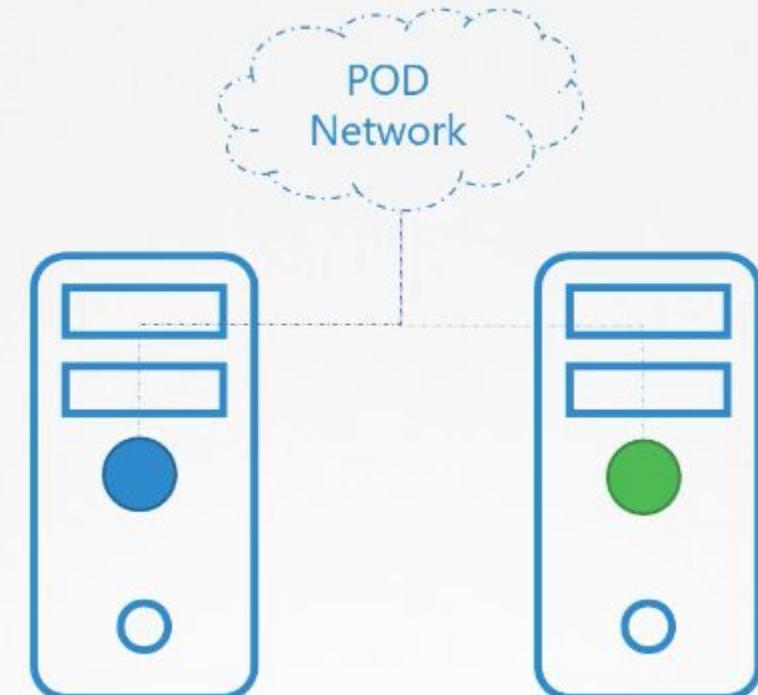
kubelet.service

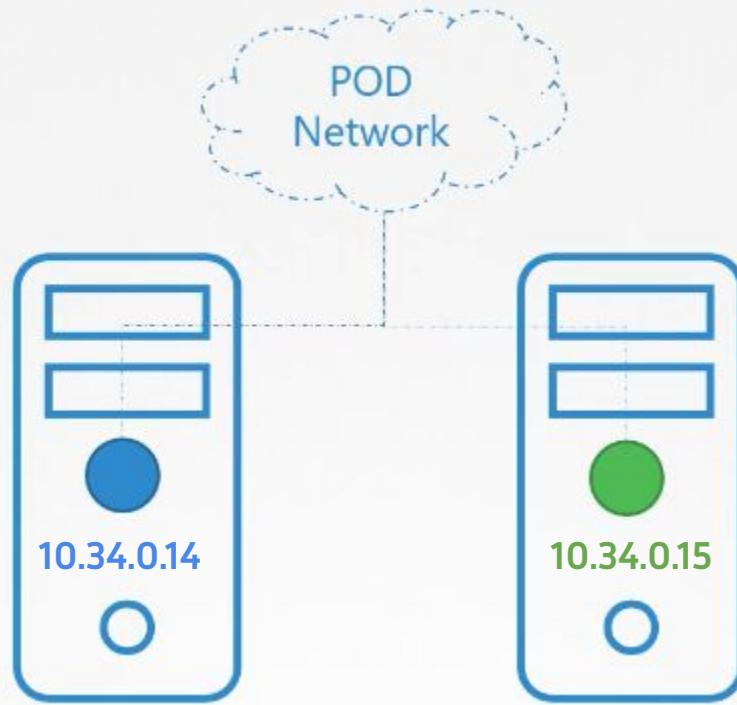
```
ExecStart=/usr/local/bin/kubelet \
--config=/var/lib/kubelet/kubelet-config.yaml \
--container-runtime=remote \
--container-runtime-endpoint=unix:///var/run/containerd/containerd.sock \
--image-pull-progress-deadline=2m \
--kubeconfig=/var/lib/kubelet/kubeconfig \
--network-plugin=cni \
--register-node=true \
--v=2
```

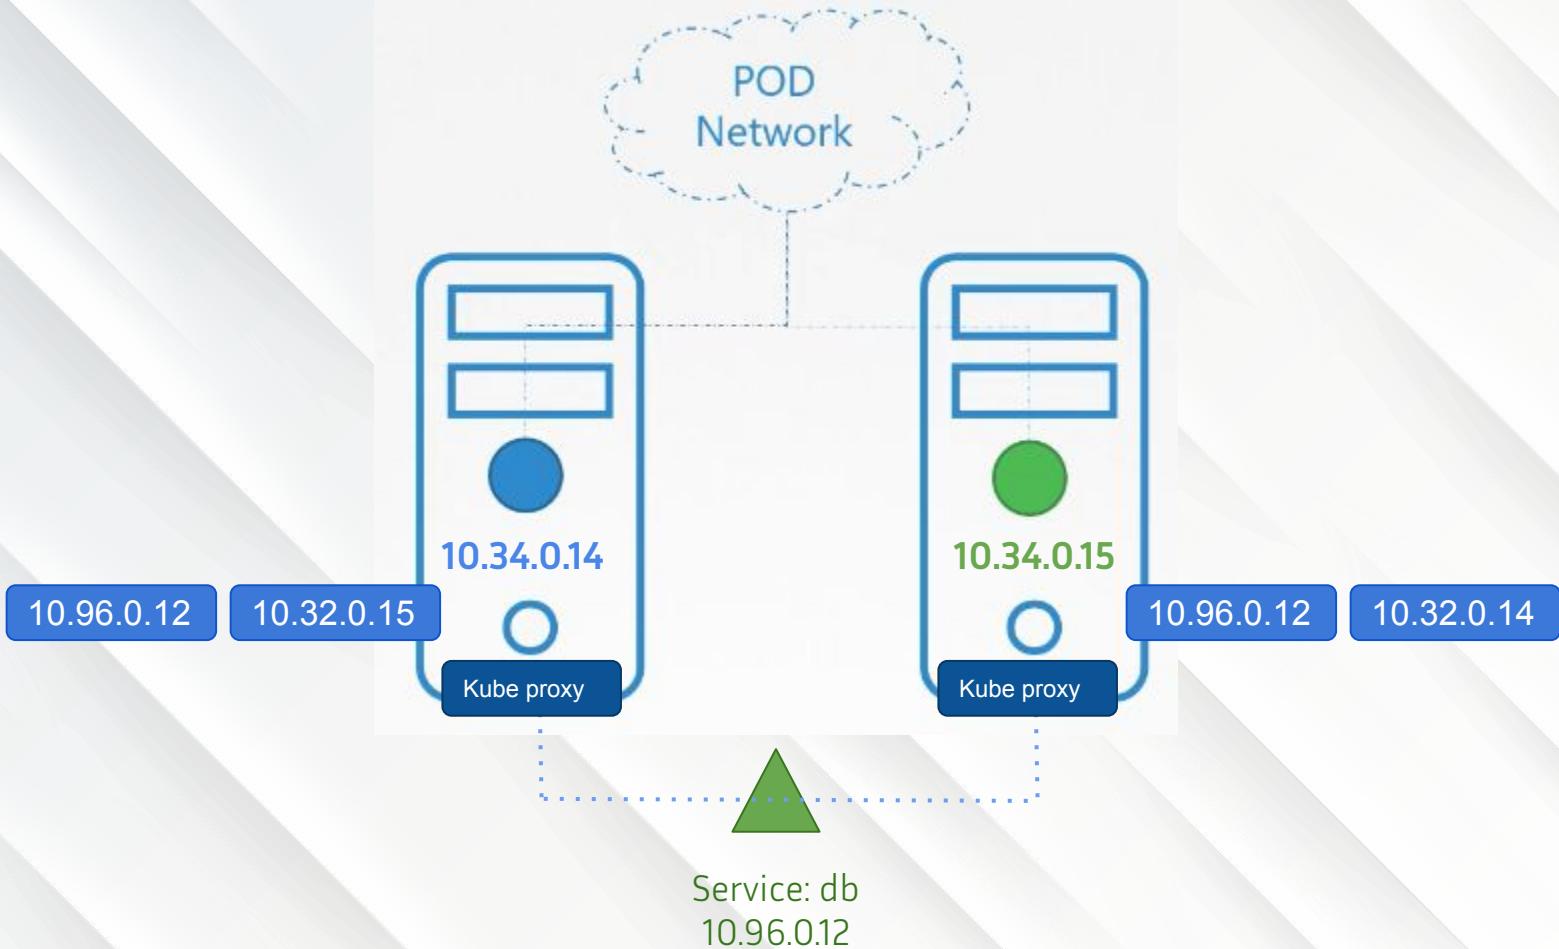
- **Docker** não possui (não precisa) do kubelet pois possui apenas um node que é o master.
- **Kubeadm** requer que você instale manualmente cada kubelet
- **Provedores** cloud automatizam para você

Kube Proxy









Instalação

```
→ kubectl get pods -n kube-system
```

NAME	READY	STATUS	RESTARTS	AGE
coredns-5644d7b6d9-2kw9p	1/1	Running	27	211d
coredns-5644d7b6d9-gdfqh	1/1	Running	27	211d
etcd-docker-desktop	1/1	Running	50	211d
kube-apiserver-docker-desktop	1/1	Running	49	211d
kube-controller-manager-docker-desktop	1/1	Running	29	211d
kube-proxy-6b9gz	1/1	Running	27	211d
kube-scheduler-docker-desktop	1/1	Running	29	211d
storage-provisioner	1/1	Running	53	211d
vpnkit-controller	1/1	Running	28	211d

```
~ took 2s
```

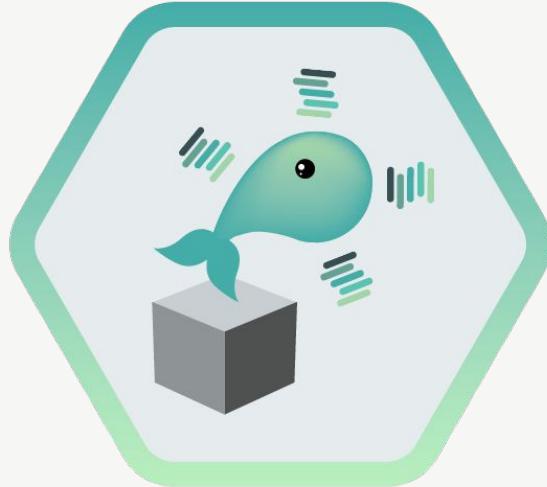
```
→ kubectl get daemonset -n kube-system
```

NAME	DESIRED	CURRENT	READY	UP-TO-DATE	AVAILABLE	NODE SELECTOR	AGE
kube-proxy	1	1	1	1	1	beta.kubernetes.io/os=	54 1h
linux	211d						

Docker

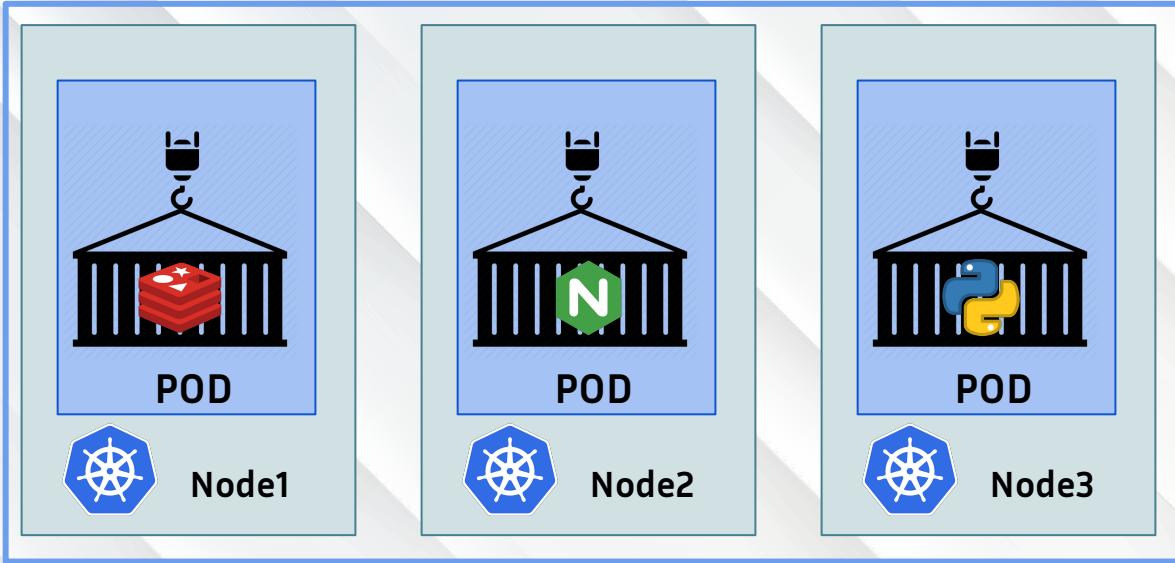


Pods



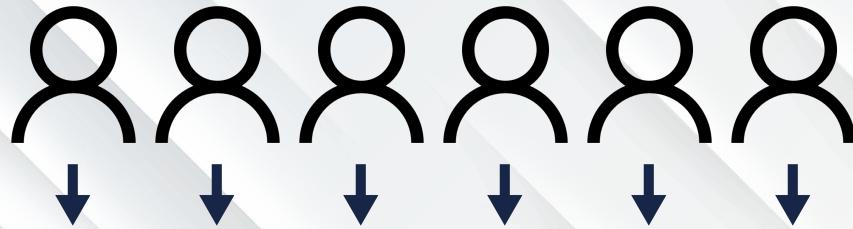
Node

d

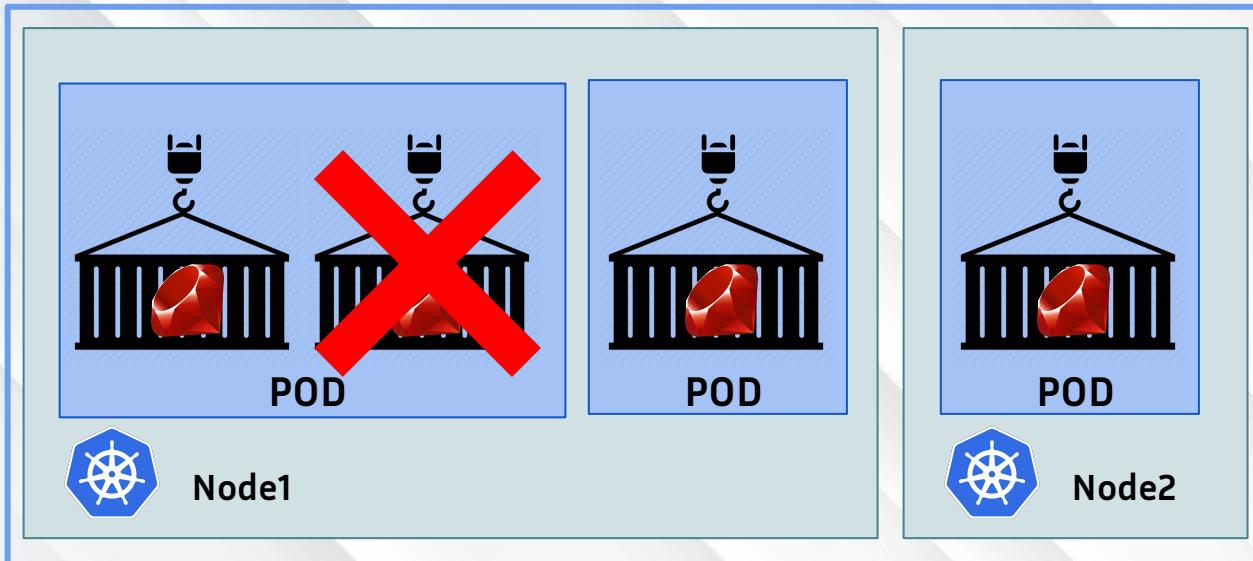


Pods

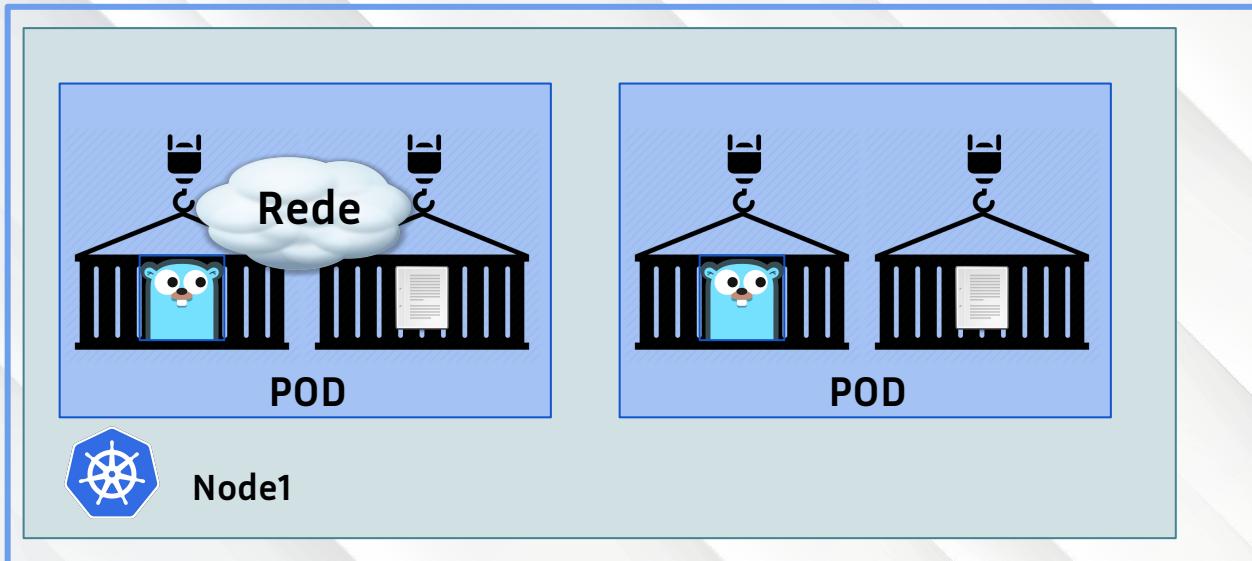




Pods



Pods - Multi-Container



Pods - kubectl

```
kubectl run nginx --image nginx
```

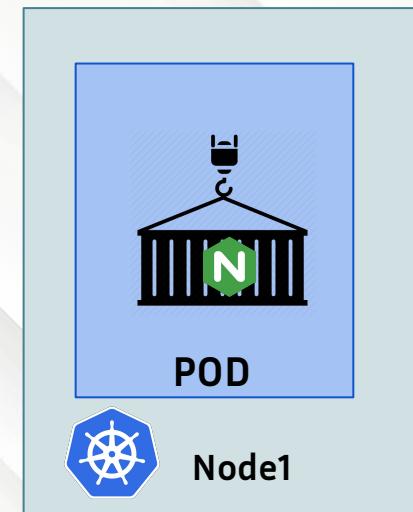
OU --image=nginx

```
kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx	0/1	ContainerCreating	0	1s

NAME	READY	STATUS	RESTARTS	AGE
nginx	1/1	Running	0	20s

```
kubectl delete pod nginx
```



Pods - kubectl - Demo & Prática

<https://www.katacoda.com/courses/kubernetes/playground>



Pods - Com YAML

```
- name: create users
  hosts: all
  tasks:
    - user:
        name: "{{ item }}"
        state: present
        groups: "{{ item }}"
      with_items:
        - { name: 'lin'
        - { name: 'lis
```



Pods - Com YAML

pod-nginx.yml

```
apiVersion: v1
kind: Pod
metadata:
  name: meuapp-pod
  labels:
    app: meuapp
    tipo: front-end
spec:
  containers:
    - name: nginx-container
      image: nginx
```

The diagram illustrates the structure of the pod-nginx.yml file. It shows the following components:

- apiVersion:** v1 (String)
- kind:** Pod (String)
- metadata:** (Dicionário)
- name:** meuapp-pod
- labels:** (Dicionário)
 - app: meuapp
 - tipo: front-end
- spec:** (Lista/Array)
 - containers:** (Dicionário)
 - name: nginx-container
 - image: nginx

Tipo de Objeto	Versão de API
POD	v1
Service	v1
ReplicaSet	apps/v1
Deployment	apps/v1

kubectl create -f pod-nginx.yml

kubectl get pods

kubectl describe pod meuapp-pod

Pods - Com YAML

```
Mounts:
  /var/run/secrets/kubernetes.io/serviceaccount from default-token-v5v8t (ro)
Conditions:
  Type        Status
  Initialized  True
  Ready       True
  ContainersReady  True
  PodScheduled  True
Volumes:
  default-token-v5v8t:
    Type:      Secret (a volume populated by a Secret)
    SecretName: default-token-v5v8t
    Optional:   false
  QoS Class:  BestEffort
  Node-Selectors: <none>
  Tolerations: node.kubernetes.io/not-ready:NoExecute for 300s
                node.kubernetes.io/unreachable:NoExecute for 300s
                4006fb9
Events:
  Type  Reason  Age   From          Message
  ----  -----  --  --  -----
  Normal Scheduled  54s  default-scheduler  Successfully assigned default/nginx to node01
  Normal Pulling   54s  kubelet, node01  Pulling image "nginx"
  Normal Pulled    48s  kubelet, node01  Successfully pulled image "nginx"
  Normal Created   47s  kubelet, node01  Created container nginx
  Normal Started   46s  kubelet, node01  Started container nginx
```

Pods (com yml)- kubectl - Demo & Prática

<https://www.katacoda.com/courses/kubernetes/playground>



kubectl describe pod <nome_do_pod>

kubectl describe pod <nome_do_pod> -o wide

kubectl delete pod <nome_do_pod>

→ kubectl run nginxsemyml --image=nginx

→ kubectl get pods

```
apiVersion: v1
kind: Pod
metadata:
  name: nginxcomyml
  labels:
    app: meunginx
    tipo: servidorweb
    autor: <seu nome>
spec:
  containers:
    - name: nginxcomyml
      image: nginx
```

Pods (com yml)- kubectl - Demo & Prática

<https://www.katacoda.com/courses/kubernetes/playground>

redispod.yml

```
apiVersion: v1
kind: Pod
metadata:
  name: rediscomyml
  labels:
    app: meuredis
    tipo: banco
    autor: <seu nome>
spec:
  containers:
    - name: rediscomyml
      image: redis
```

kubectl create -f <nome_do_arquivo>

kubectl delete pods -l tipo=servidorweb

apachepod.yml

```
apiVersion: v1
kind: Pod
metadata:
  name: apachecomyml
  labels:
    app: meuapache
    tipo: servidorweb
    autor: <seu nome>
spec:
  containers:
    - name: apachecomyml
      image: httpd
```

ou kubectl create -f nginxpod.yml -f apachepod.yml
 -f redispod.yml

Pods (com yml)- kubectl - Demo & Prática

<https://www.katacoda.com/courses/kubernetes/playground>

Arquivo: **pod2containers.yml**

Contaiier 1:

- **Image:** nginx

Container 2:

- **Image:** fluentd

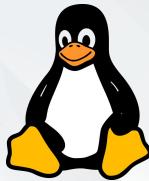
```
apiVersion: v1
kind: Pod
metadata:
  name: rediscomyml
  labels:
    app: meureredis
    tipo: banco
    autor: <seu nome>
spec:
  containers:
    - name: rediscomyml
      image: redis
    - name: algum-nome
      image: alguma-imagem
```



Pods - Quizzzzzz



```
brew install --cask google-cloud-sdk
```



```
curl -O  
https://dl.google.com/dl/cloudsdk/channels/rapid/downlo  
ads/google-cloud-sdk-323.0.0-linux-x86_64.tar.gz
```

```
tar -xf google-cloud-sdk-323.0.0-linux-x86_64.tar.gz
```

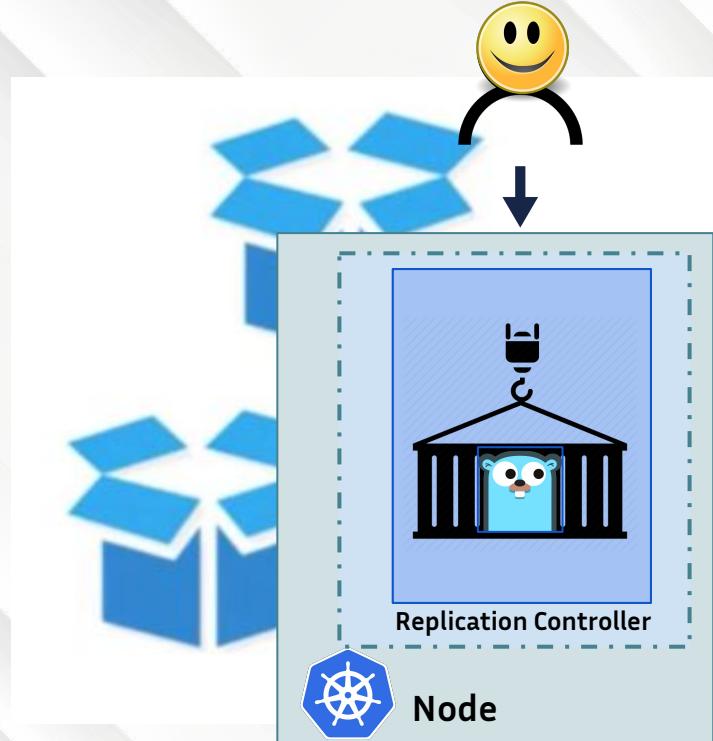
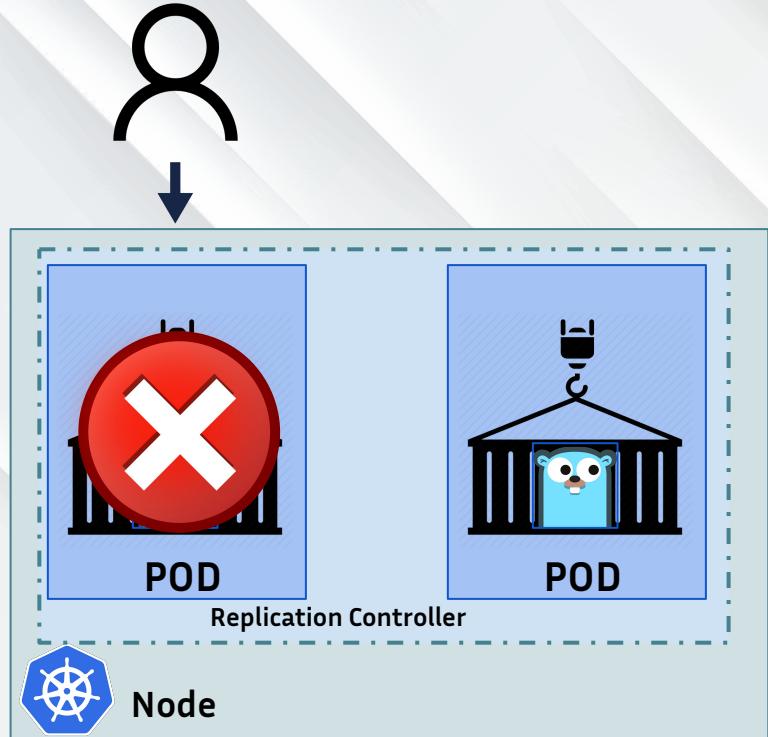
```
./google-cloud-sdk/install.sh
```



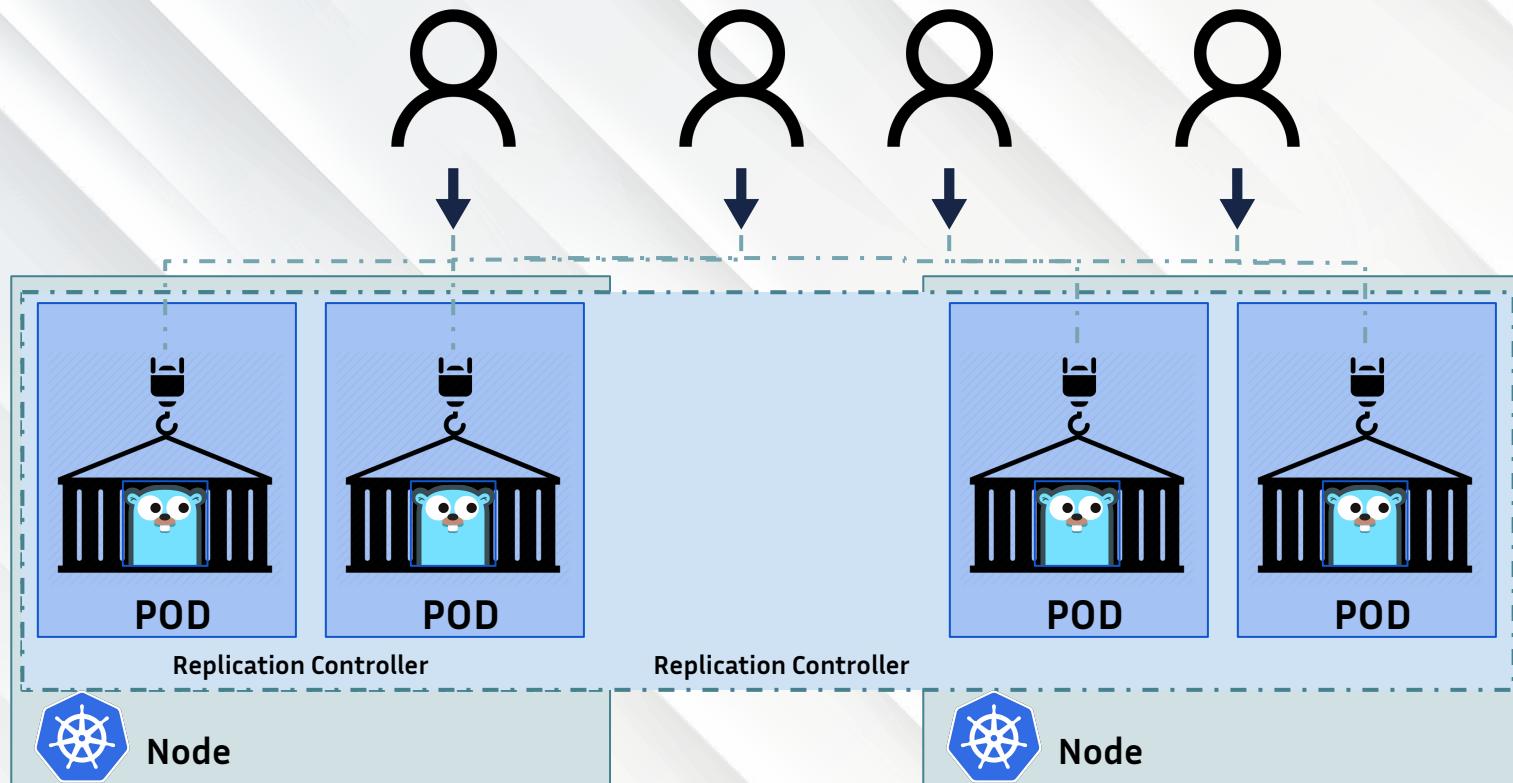
```
(New-Object  
Net.WebClient).DownloadFile("https://dl.google.com/  
dl/cloudsdk/channels/rapid/GoogleCloudSDKInstaller.  
exe", "$env:Temp\GoogleCloudSDKInstaller.exe")  
  
& $env:Temp\GoogleCloudSDKInstaller.exe
```

dsdk/channels/rapid/Go

Kubernetes Controllers - Replication Controller



Load Balancing & Scaling



Replication Controller & Replica Set

Replication Controller

Replica Set

Replication Controller

rc-teste.yml

```
apiVersion: v1
kind: ReplicationController
metadata:
  name: meuapp-rc
  labels:
    app: meuapp
    tipo: front-end
spec:
  template:
    metadata:
      name: meuapp-pod
      labels:
        app: meuapp
        tipo: front-end
    spec:
      containers:
        - name: nginx-container
          image: nginx
    replicas: 3
```

pod-teste.yml

```
apiVersion: v1
kind: Pod
```

```
> kubectl create -f rc-teste.yml
```

```
Replicationcontroller "meuapp-rc" created
```

Replication Controller

```
> kubectl get replicationcontroller
```

NAME	DESIRED	CURRENT	READY	AGE
Meuapp-rc	3	3	3	19s

```
> kubectl get pods
```

NAME	DESIRED	CURRENT	READY	AGE
Meuapp-rc-41vk9	3	3	3	19s
Meuapp-rc-23y4h	3	3	3	19s
Meuapp-rc-4hgv7	3	3	3	19s

Replication Controller - Prática

```
> gcloud container clusters get-credentials possible-sun-222819-gke --region us-central1
```

<seunome>-app-rc.yml

```
apiVersion: v1
kind: ReplicationController
metadata:
  name: <nome>-app-rc
  labels:
    app: <nome>-app
    tipo: web
spec:
  template:
    metadata:
      name: <nome>-app-pod
      labels:
        app: <nome>-app
        tipo: web
    spec:
      containers:
        - name: <nome>-nginx-container
          image: nginx
replicas: 3
```

```
> kubectl create -f <seunome>-app-rc.yml
```

```
> kubectl get replicationcontroller
```



ReplicaSet

meuapp-replicaset.yml

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: meuapp-replicaset
  labels:
    app: meuapp
    tipo: front-end
spec:
  template:
    metadata:
      name: meuapp-pod
      labels:
        app: meuapp
        tipo: front-end
    spec:
      containers:
      - name: nginx-container
        image: nginx
replicas: 4
```

meupod.yml

```
apiVersion: v1
kind: Pod
```

ReplicaSet

```
> kubectl create -f  
meuapp-replicaset.yml
```

```
Replicaset "meuapp-replicaset" created
```

```
> kubectl get replicaset
```

NAME	DESIRED	CURRENT	READY	AGE
replicaset	3	3	3	19s

```
> kubectl get pods
```

NAME	DESIRED	CURRENT	READY	AGE
repset-41v9	3	3	3	19s
repaset-234h	3	3	3	19s
recaset-4hv7	3	3	3	19s

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: meuapp-replicaset
  labels:
    app: meuapp
    tipo: front-end
spec:
  template:
    metadata:
      name: meuapp-pod
      labels:
        app: meuapp
        tipo: front-end
    spec:
      containers:
        - name: nginx-container
          image: nginx
      replicas: 3
      selector:
        matchLabels:
          type: front-end
```



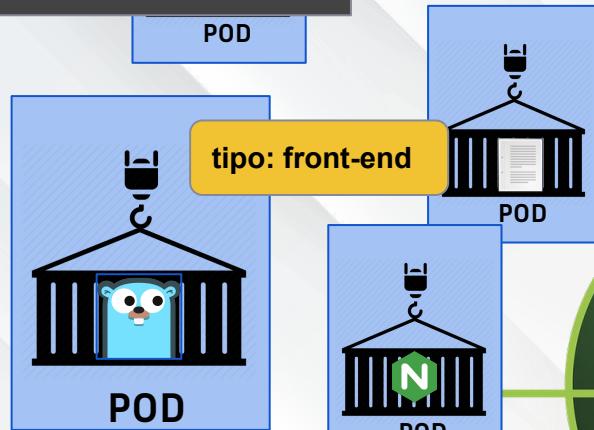
Labels and Selectors

replicaset.yml

```
selector:  
  matchLabels:  
    tipo: front-end
```



POD



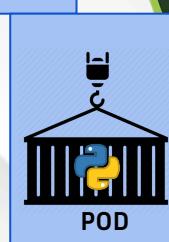
POD



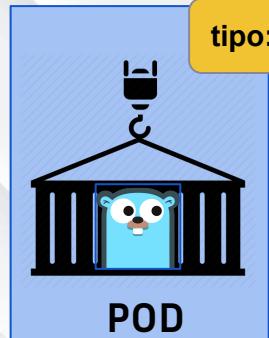
POD



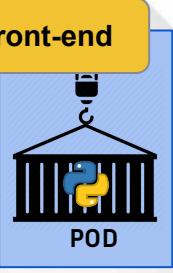
POD



POD



tipo: front-end



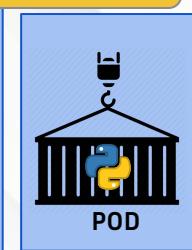
POD



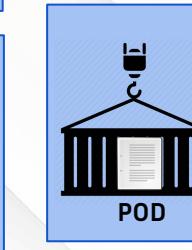
POD



POD



POD



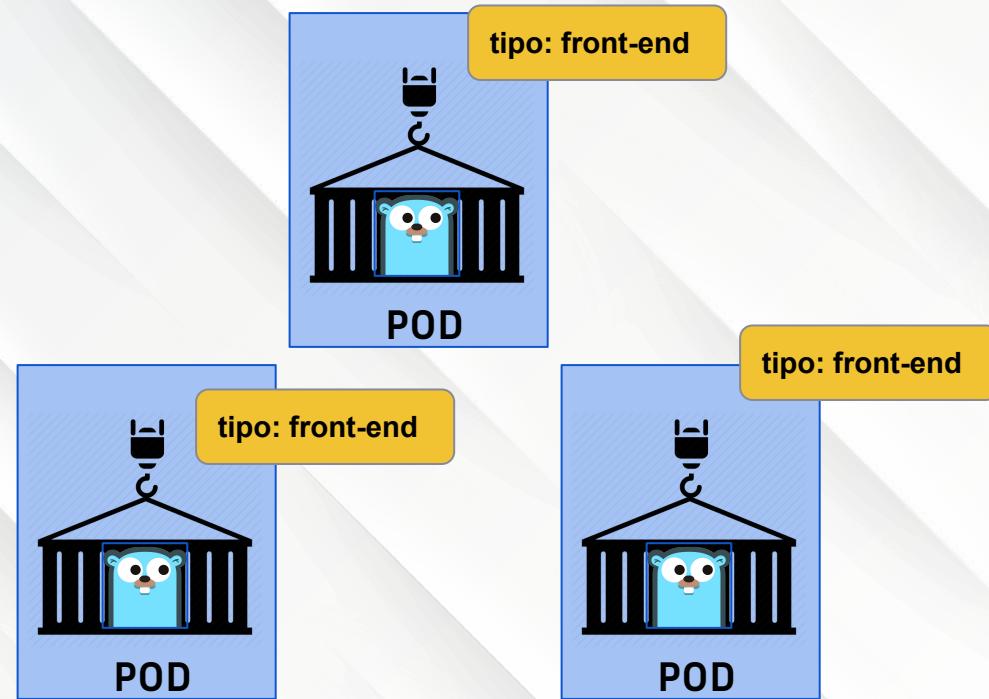
POD

metadata:
 name: meuapp-pod
 labels:
 tipo: front-end



replicaset.yml

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: meuapp-replicaset
  labels:
    app: meuapp
    type: front-end
spec:
  template:
    metadata:
      name: meu-app-pod
      labels:
        app: meu-app
        tipo: front-end
    spec:
      containers:
        - name: nginx-container
          image: nginx
  replicas: 3
  selector:
    matchLabels:
      tipo: front-end
```



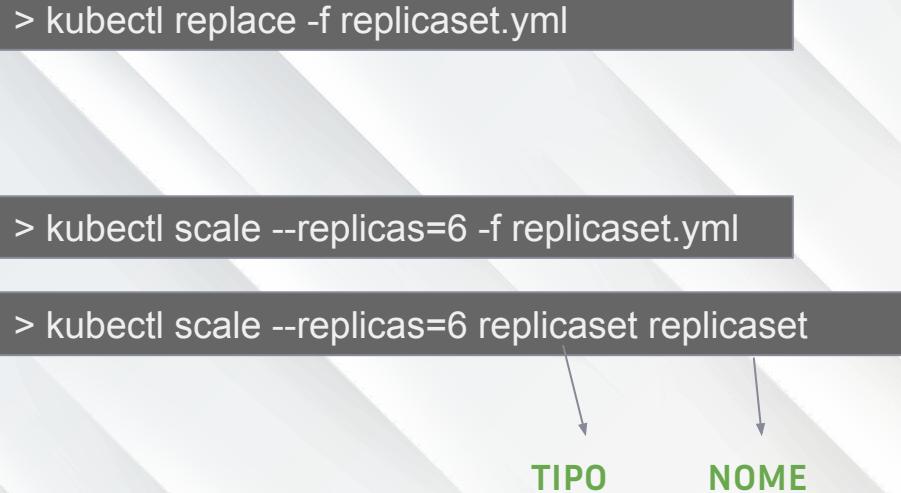
Scale

```
> kubectl replace -f replicaset.yml
```

```
> kubectl scale --replicas=6 -f replicaset.yml
```

```
> kubectl scale --replicas=6 replicaset replicaset
```

TIPO NOME



The diagram consists of three separate command examples from the left side, each enclosed in a dark grey box. Arrows from each example point downwards towards the 'TIPO' and 'NOME' labels located in the center of the slide.

replicaset.yml

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: meuapp-replicaset
  labels:
    app: meuapp
    type: front-end
spec:
  template:
    metadata:
      name: meu-app-pod
      labels:
        app: meu-app
        tipo: front-end
    spec:
      containers:
        - name: nginx-container
          image: nginx
  replicas: 6
  selector:
    matchLabels:
      tipo: front-end
```

ReplicaSet - Prática

Conectando-se ao cluster:

```
> gcloud container clusters get-credentials possible-sun-222819-gke --region us-central1
```

Testando o cluster:

```
> kubectl get nodes
```



```
<nome>-replicaset.yml
```

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: <nome>-replicaset
  labels:
    app: <nome>-meuapp
    type: <nome>-web
spec:
  template:
    metadata:
      name: <nome>-app-pod
      labels:
        app: <nome>-app
        tipo: <nome>-web
    spec:
      containers:
        - name: nginx-container
          image: nginx
  replicas: 3
  selector:
    matchLabels:
      tipo: <nome>-web
```

Criando o replica set: **ReplicaSet - Prática**

```
> kubectl create -f <nome>-replicaset.yml
```

Listando replicaset:

```
> kubectl get replicaset
```

Listando pods:

```
> kubectl get pods
```

Listando tudo:

```
> kubectl get all
```

Outra opção:

```
> kubectl get rs
```

```
> kubectl replace -f <nome>-replicaset.yml
```

```
> kubectl scale --replicas=6 -f <nome>-replicaset.yml
```

```
> kubectl scale --replicas=6 replicaset <noem>-replicaset.yml
```

<nome>-pod.yml

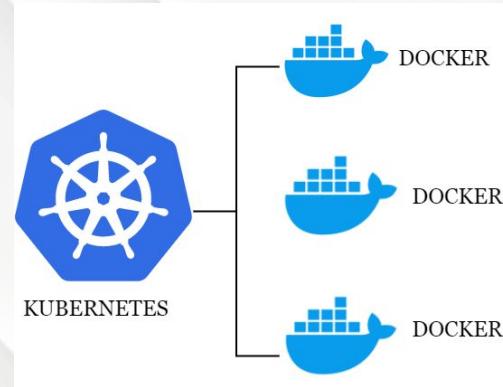
```
apiVersion: v1
kind: Pod
metadata:
  name: <nome>-app-pod
  labels:
    app: <nome>-app
    tipo: <nome>-web
    autor: <seu nome>
spec:
  containers:
    - name: nginx-container
      image: nginx
```



Kind - Instalação

shorturl.at/avMZ3

Kubernetes in Docker



Kind - Setup

Listar containers rodando no sistema:

```
> docker ps
```

Criando o cluster:

```
> kind create cluster --name teste
```

Verifiquem informações do cluster:

```
> kubectl cluster-info --context kind-teste
```

Verifiquem se estão no contexto correto:

```
> kubectl config get-contexts
```

Listar clusters:

```
> kind get clusters
```

Deletar clusters:

```
> kind delete cluster --name teste
```

Kind - Setup

Definições do cluster:

```
> cluster-teste.yml
```

```
kind: Cluster
apiVersion: kind.x-k8s.io/v1alpha4
nodes:
- role: control-plane
- role: worker
- role: worker
```

Criando o cluster:

```
> kind create cluster --name teste --config cluster-teste.yml
```

Listar containers rodando no sistema:

```
> docker ps
```

Verifiquem informações do cluster:

```
> kubectl cluster-info --context kind-teste
```

```
$ kind create cluster --name teste --config cluster-teste.yml
Creating cluster "teste" ...
✓ Ensuring node image (kindest/node:v1.19.1) 
✓ Preparing nodes 
✓ Writing configuration 
✓ Starting control-plane 
✓ Installing CNI 
✓ Installing StorageClass 
✓ Joining worker nodes 
Set kubectl context to "kind-teste"
You can now use your cluster with:

kubectl cluster-info --context kind-teste

Have a nice day! 
```

Listar clusters:

```
> kind get clusters
```

Kind - Prática

nginx-deployment.yml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: nginx-deployment
spec:
  selector:
    matchLabels:
      app: nginx
  replicas: 2
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
          image: nginx:1.16.1
          ports:
            - containerPort: 80
```

nginx-service.yml

```
apiVersion: v1
kind: Service
metadata:
  name: nginx-service
  labels:
    app: nginx
spec:
  type: NodePort
  ports:
    - port: 80
      nodePort: 30225
  selector:
    app: nginx
```

Kind - Prática

Criando o deployment:

```
> kubectl create -f nginx-deployment.yml
```

Criando o service:

```
> kubectl create -f nginx-service.yml
```

Vendo tudo:

```
> kubectl get all
```

127.0.0.1:8000

Mapeando porta:

```
> kubectl port-forward service/nginx-service 8000:80
```

Forwarding from 127.0.0.1:8000 -> 80

Forwarding from [::1]:8000 -> 80

Handling connection for 8000

Handling connection for 8000

Welcome to nginx!

If you see this page, the nginx web server is successfully installed and working. Further configuration is required.

For online documentation and support please refer to nginx.org.
Commercial support is available at nginx.com.

Thank you for using nginx.

Deployments

